

MS&E 233

Game Theory, Data Science and AI

Lecture 8

Vasilis Syrgkanis

Assistant Professor

Management Science and Engineering

(by courtesy) Computer Science and Electrical Engineering

Institute for Computational and Mathematical Engineering

Computational Game Theory for Complex Games

- Basics of game theory and zero-sum games (T)
- Basics of online learning theory (T)
- Solving zero-sum games via online learning (T)
- 1 • *HW1: implement simple algorithms to solve zero-sum games*
- Applications to ML and AI (T+A)
- *HW2: implement boosting as solving a zero-sum game*

- Basics of extensive-form games
- Solving extensive-form games via online learning (T)
- 2 • *HW3: implement agents to solve very simple variants of poker*

- General games, equilibria and online learning (T)
- 3 • **Online learning in general games**
- *HW4: implement no-regret algorithms that converge to correlated equilibria in general games*

Data Science for Auctions and Mechanisms

- Basics and applications of auction theory (T+A)
- 4 • **Learning to bid in auctions via online learning (T)**
- *HW5: implement bandit algorithms to bid in ad auctions*

- Optimal auctions and mechanisms (T)
- Simple vs optimal mechanisms (T)
- 5 • *HW6: calculate equilibria in simple auctions, implement simple and optimal auctions, analyze revenue empirically*

- Optimizing mechanisms from samples (T)
- Online optimization of auctions and mechanisms (T)
- 6 • *HW7: implement procedures to learn approximately optimal auctions from historical samples and in an online manner*

Further Topics

- Econometrics in games and auctions (T+A)
- A/B testing in markets (T+A)
- 7 • *HW8: implement procedure to estimate values from bids in an auction, empirically analyze inaccuracy of A/B tests in markets*

Guest Lectures

- Mechanism Design for LLMs, Renato Paes Leme, Google Research
- Auto-bidding in Sponsored Search Auctions, Kshipra Bhawalkar, Google Research

Recap: Regret vs Correlated Equilibrium

- No-regret property, implies

Distributions that satisfy this are called **Coarse Correlated Equilibria**

$$\forall s'_i: \sum_s \pi^T(s) \left(u_i(s) - u_i(s'_i, s_{-i}) \right) \geq -\tilde{\epsilon}(T, \delta) \rightarrow 0$$

- Correlated equilibrium requires conditioning on recommendation

$$\forall s_i^*, s'_i: \sum_{s: s_i = s_i^*} \pi^T(s) \left(u_i(s) - u_i(s'_i, s_{-i}) \right) \geq 0$$

At subset of periods
when **played** s_i^*



You don't regret
switching to s'_i


Recap: Swaps and Correlated Equilibrium

- Correlated equilibrium requires conditioning on recommendation

$$\forall s_i^*, s'_i: \sum_{s: s_i = s_i^*} \pi^T(s) \left(u_i(s) - u_i(s'_i, s_{-i}) \right) \geq 0$$

- Equivalently: for any **swap** function ϕ that maps original actions s_i to deviating actions s'_i (potentially different for each original s_i)

$$\sum_s \pi^T(s) \left(u_i(s) - u_i(\phi(s_i), s_{-i}) \right) \geq 0$$



Recap: No-Swap Regret!

- No-regret property requires

$$\frac{1}{T} \sum_{t=1}^T u_i(s^t) \geq \max_{s'_i \in S_i} \frac{1}{T} \sum_{t=1}^T u_i(s'_i, s_{-i}^t) - \tilde{\epsilon}(T, \delta)$$

- No-swap regret property requires

$$\forall \phi: \frac{1}{T} \sum_{t=1}^T u_i(s^t) \geq \frac{1}{T} \sum_{t=1}^T u_i(\phi(s_i^t), s_{-i}^t) - \tilde{\epsilon}(T, \delta)$$

Theorem. If all players use no-swap regret algorithms, then the empirical joint distribution converges to a Correlated Equilibrium


Can we construct algorithms with
vanishing no-swap regret?

No Swap Regret vs No Regret

- At period t you choose action i_t from distribution x_t over n actions
- Observe vector $\ell_t = (\ell_t^1, \dots, \ell_t^n)$ containing loss of each action
- You incur the loss of the action you chose $\ell_t^{i_t}$
- No-regret: for any action i , you do not regret always taking action i

$$\frac{1}{T} \sum_t \ell_t^{i_t} \leq \frac{1}{T} \sum_t \ell_t^i + \tilde{\epsilon}(T, \delta), \quad \text{w. p. } 1 - \delta$$

No-Regret




No Swap Regret vs No Regret


- At period t you choose action i_t from distribution x_t over n actions
- Observe vector $\ell_t = (\ell_t^1, \dots, \ell_t^n)$ containing loss of each action
- You incur the loss of the action you chose $\ell_t^{i_t}$
- No-swap regret: for any swap function ϕ mapping original actions i to alternatives $i' = \phi(i)$, you do not regret making that swap

$$\frac{1}{T} \sum_t \ell_t^{i_t} \leq \frac{1}{T} \sum_t \ell_t^{\phi(i_t)} + \tilde{\epsilon}(T, \delta), \quad \text{w. p. } 1 - \delta$$


No-Swap Regret



Alternatives




No-Swap Regret



Alternatives




No-Swap Regret



Alternatives

Switch to 1
when playing 2

Swap
1 → 1
2 → 1
3 → 3




Switch to 3
when playing 2

Swap
1 → 1
2 → 3
3 → 3



No-Swap Regret



Vanishing regret for complex swaps is implied by vanishing regret of simple swaps:
switch to j' whenever you had played j and leave everything else as is

No Swap Regret vs No Regret

- **No-swap regret:** for any swap function ϕ mapping original actions i to alternatives $i' = \phi(i)$, you do not regret making that swap

$$\frac{1}{T} \sum_t \ell_t^{i_t} \leq \frac{1}{T} \sum_t \ell_t^{\phi(i_t)} + \tilde{\epsilon}(T, \delta), \quad \text{w. p. } 1 - \delta$$

- **Equivalently:** for **subset of periods** when you played i you don't regret any other action i'

$$\frac{1}{T} \sum_{t: i_t = i} \ell_t^{i_t} \leq \max_{i'} \frac{1}{T} \sum_{t: i_t = i} \ell_t^{i'} + \tilde{\epsilon}(T, \delta), \quad \text{w. p. } 1 - \delta$$

You have an online learning problem, for simplicity, with 2 actions. Is any no-swap regret sequence a no-regret sequence?

Yes

0%

No

0%

You have an online learning problem, for simplicity, with 2 actions. Is any no-regret sequence a no-swap regret sequence?

Yes

0%

No

0%

You have an online learning problem, for simplicity, with 3 actions. Is any no-regret sequence a no-swap regret sequence?

Yes

0%


No

0%

No-Swap Regret

Action 1

Action 2



Alternatives

Switch to 1
when playing 2

Swap

1→1
2→1



Total Loss = 5

Switch to 2
when playing 1


Swap

1→2
2→2



Total Loss = 6

No-Swap Regret



Alternatives

Switch to 1
when playing 2

Swap 1→1
2→1




No-swap regret is weirdly implied by no-regret when you only have two actions.

Intuition: no-regret towards action j is the same as no-regret on the subset of periods when you did not play j . With two actions, these are exactly the periods when you played j'

Switch to 2
when playing 1

Swap 1→2
2→2



Can we reduce no-swap regret to
no-regret?

No Swap Regret vs No Regret

- **For subset of periods** when played i don't regret any other i'

$$\frac{1}{T} \sum_{t: i_t = i} \ell_t^{i_t} \leq \max_{i'} \frac{1}{T} \sum_{t: i_t = i} \ell_t^{i'} + \tilde{\epsilon}(T, \delta), \quad \text{w. p. } 1 - \delta$$

- This looks like the no-regret property, but on a subset of periods
- If ahead of time we knew on which subset of periods we'd play i
- We could spawn a separate no-regret algorithm A_i
- When it was time to play i we would call A_i and report back loss

Swap to No-Regret Reduction

actions

1

⋮

j

⋮

n

Master Algorithm (M)

A_1

Responsible for controlling regret
in periods when 1 was played

⋮

A_i

Responsible for controlling regret
in periods when i was played

⋮

A_n

Responsible for controlling regret
in periods when n was played

Swap to No-Regret Reduction

actions


1

⋮

j

⋮

n



1 Choose algorithm i_t based on probability distribution q_t

A_1
Responsible for controlling regret
in periods when 1 was played q_t^1

⋮

A_i
Responsible for controlling regret
in periods when i was played q_t^i

⋮

A_n
Responsible for controlling regret
in periods when n was played q_t^n

Swap to No-Regret Reduction

actions

1

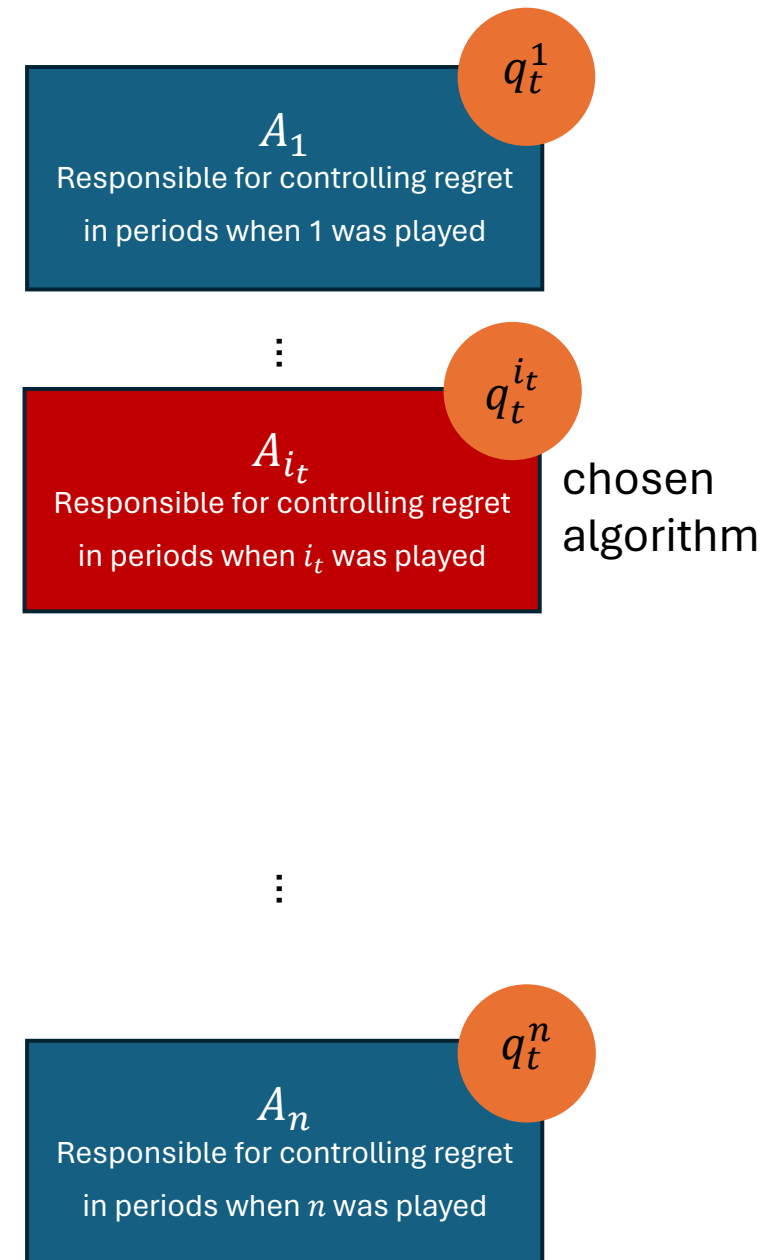
\vdots

j

\vdots

n

Master Algorithm (M)



Swap to No-Regret Reduction

actions

1

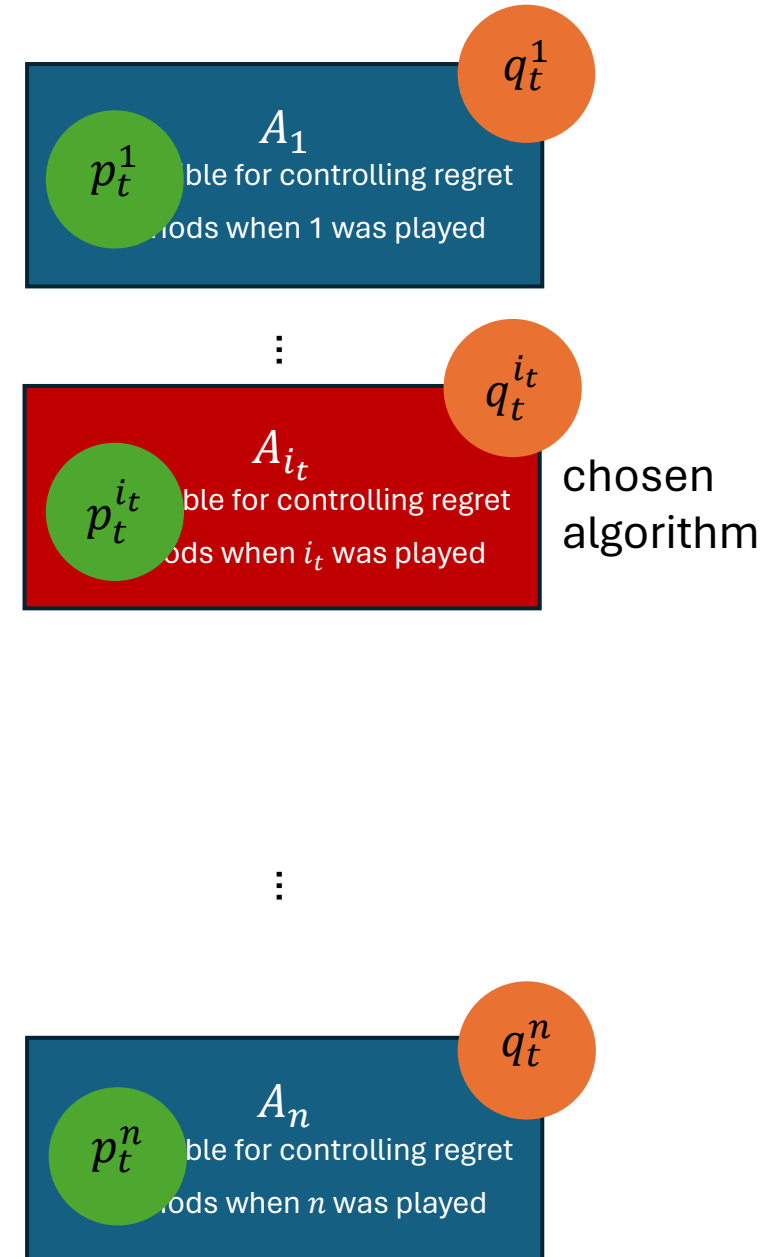
⋮

j

⋮

n

Master Algorithm (M)



Swap to No-Regret Reduction

actions


1

⋮

j

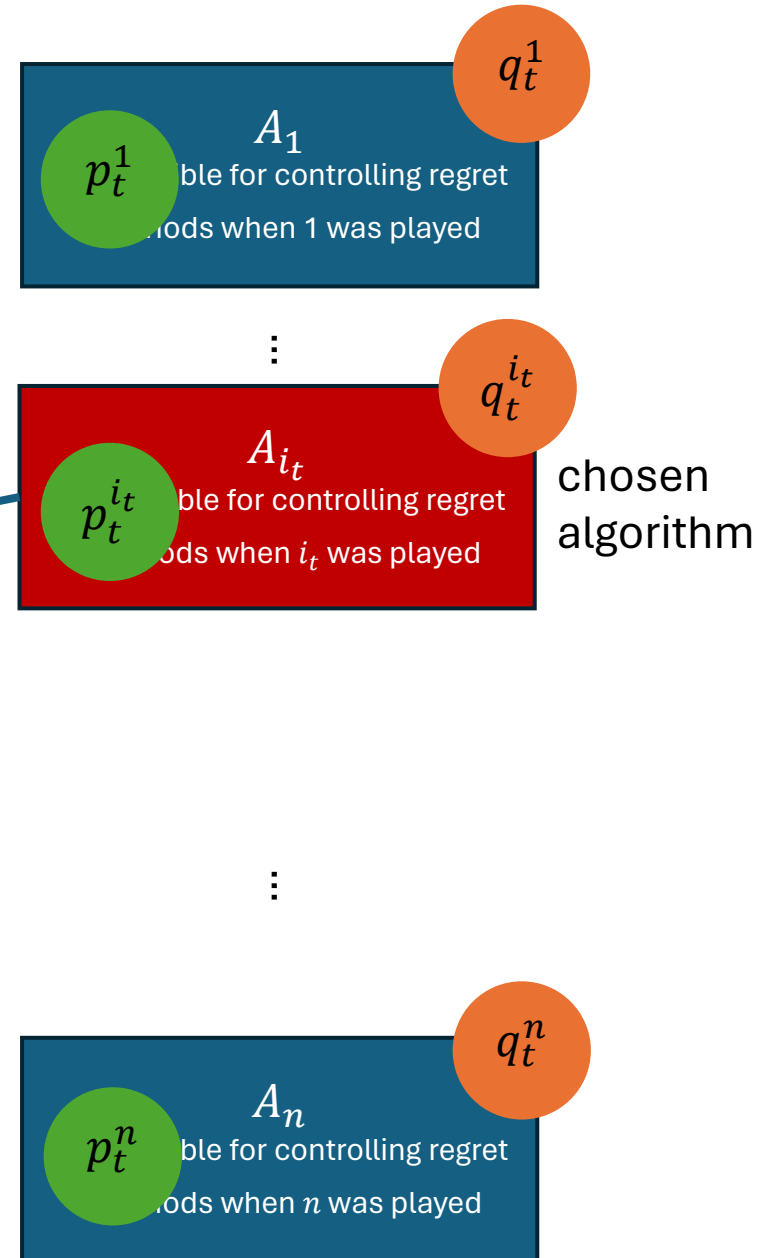
⋮

n

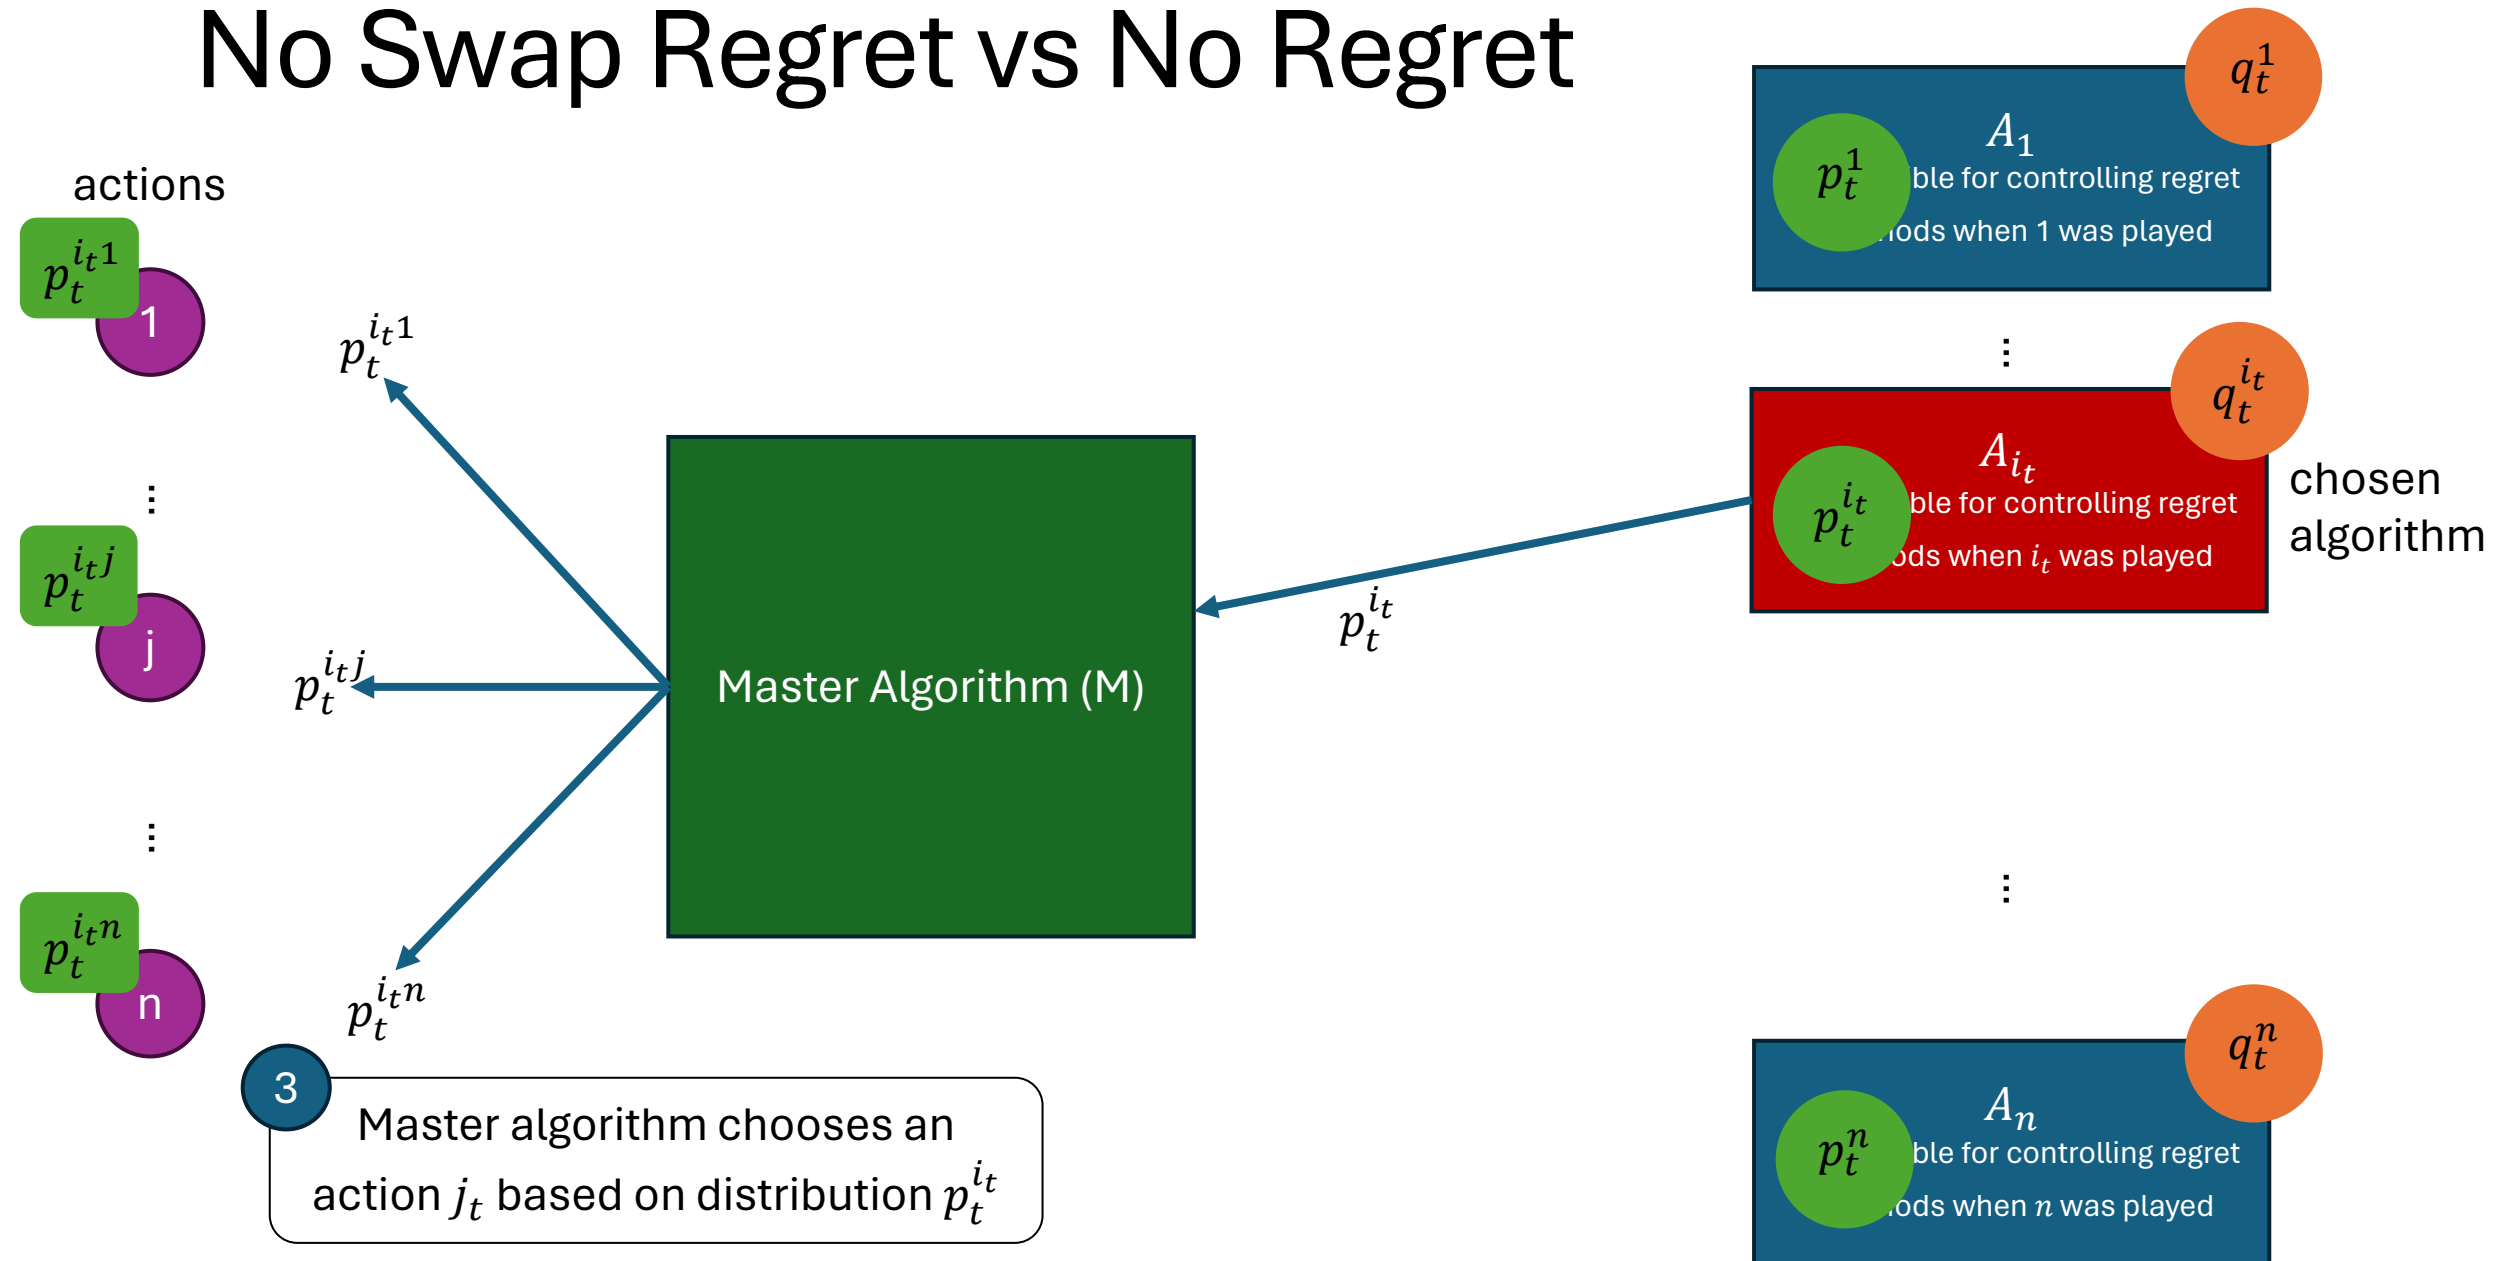


2
Algorithm A_{i_t} reports
some probability
distribution $p_t^{i_t}$ over
actions


$p_t^{i_t}$



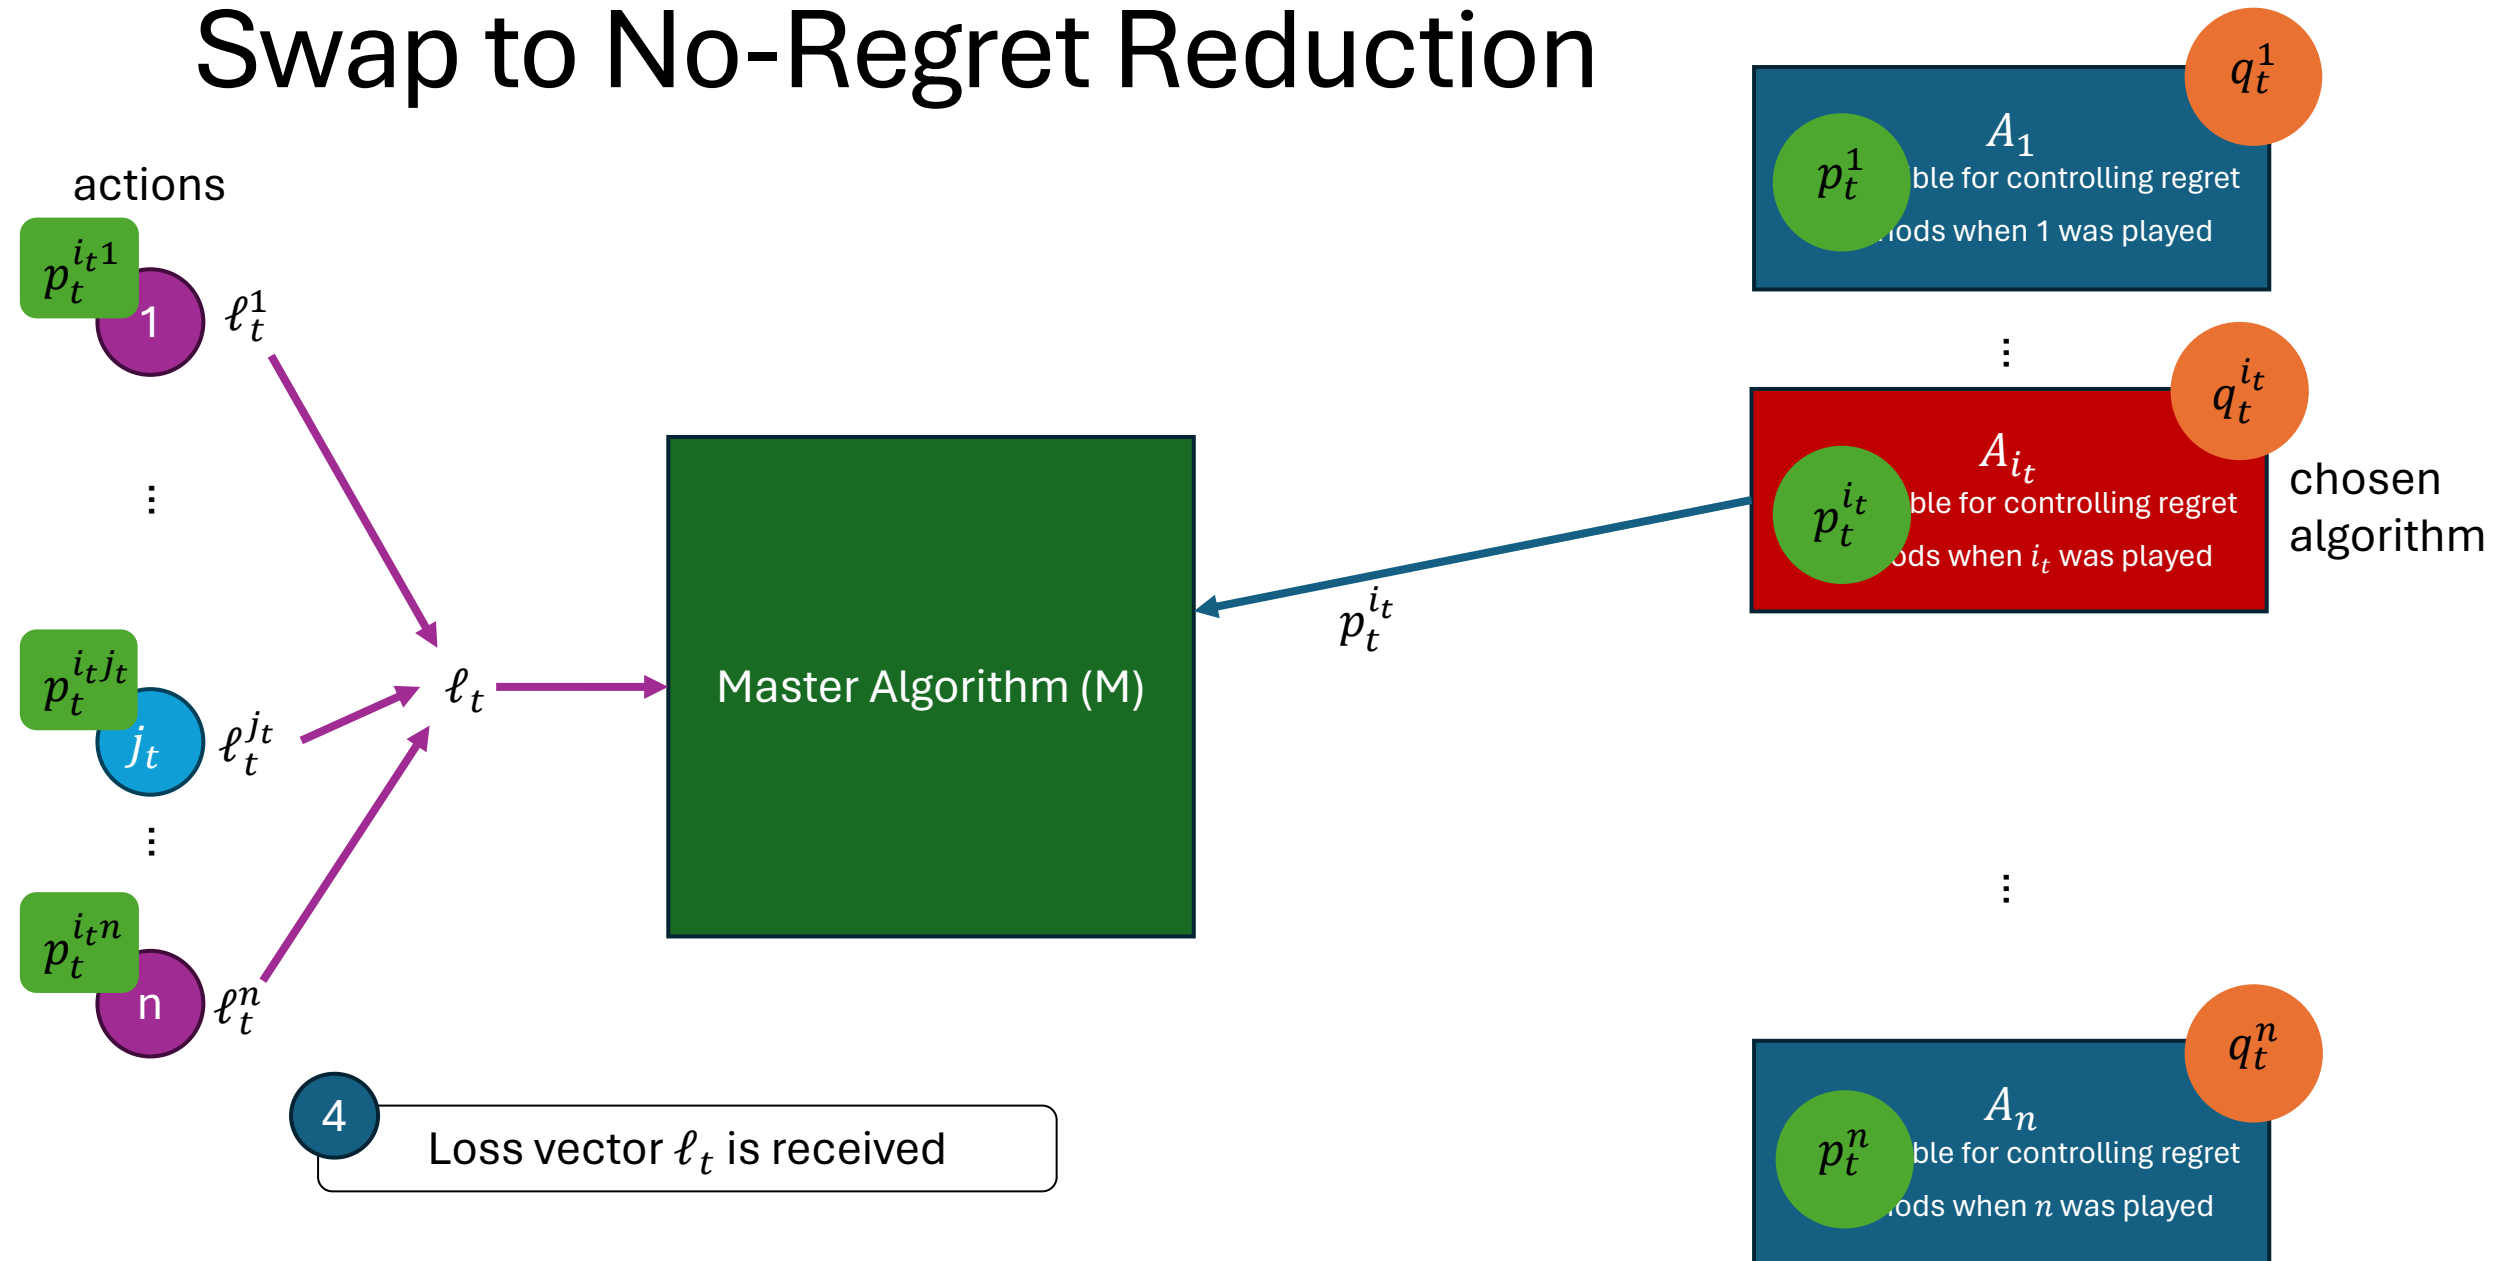
No Swap Regret vs No Regret



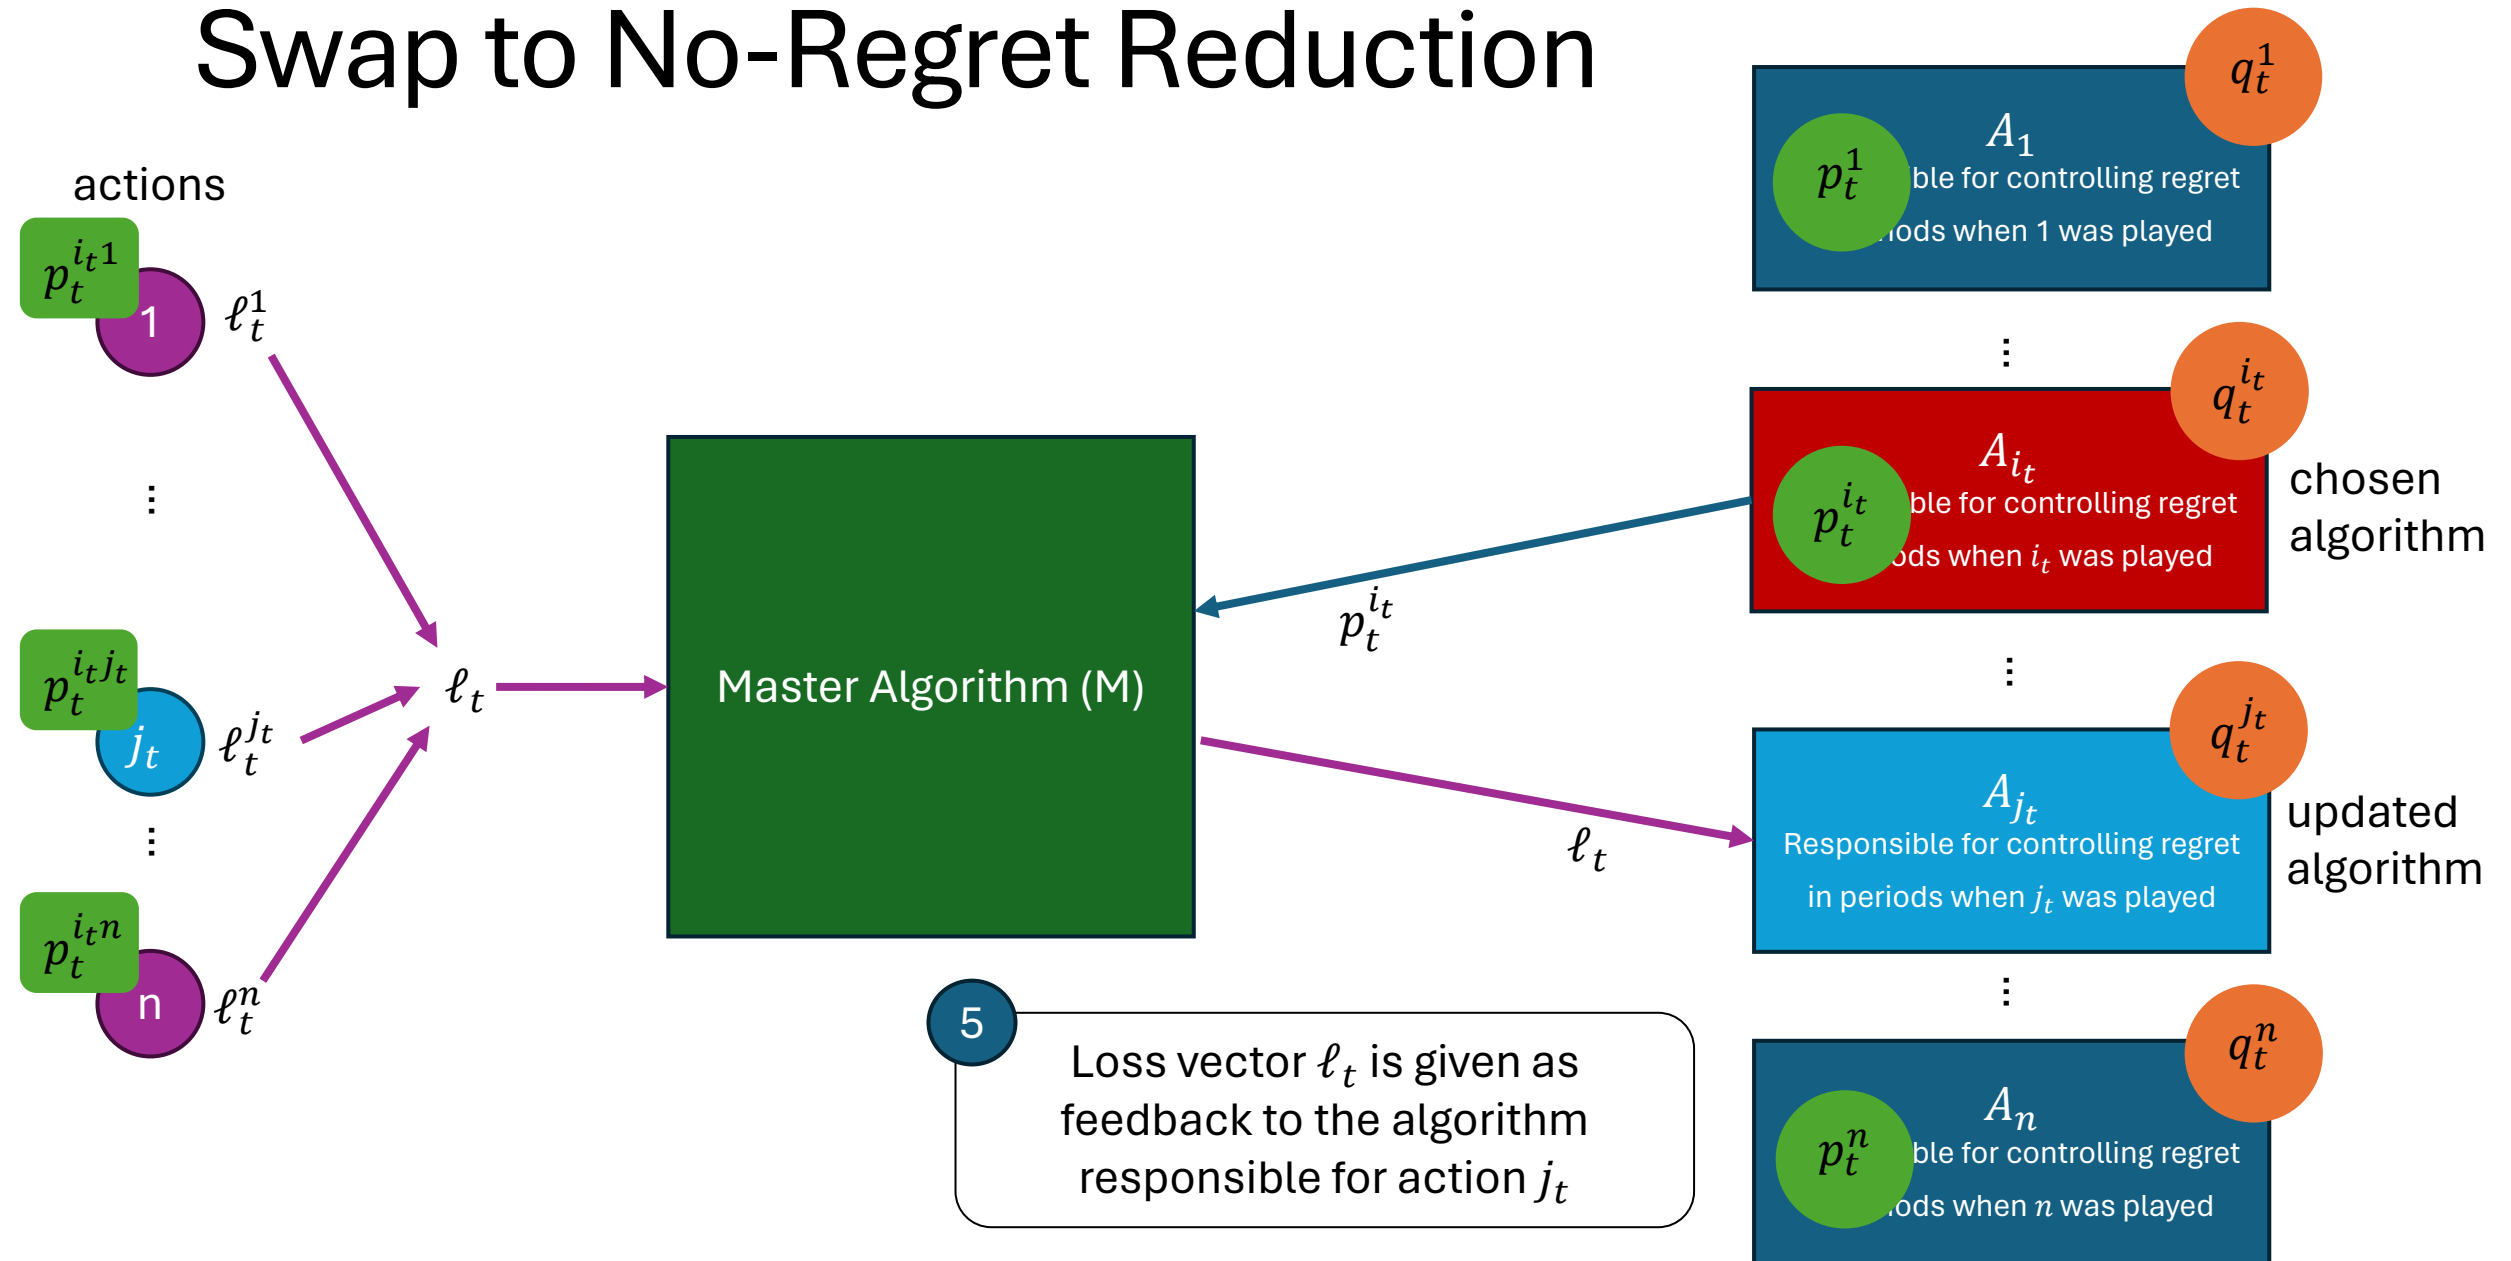
Swap to No-Regret Reduction



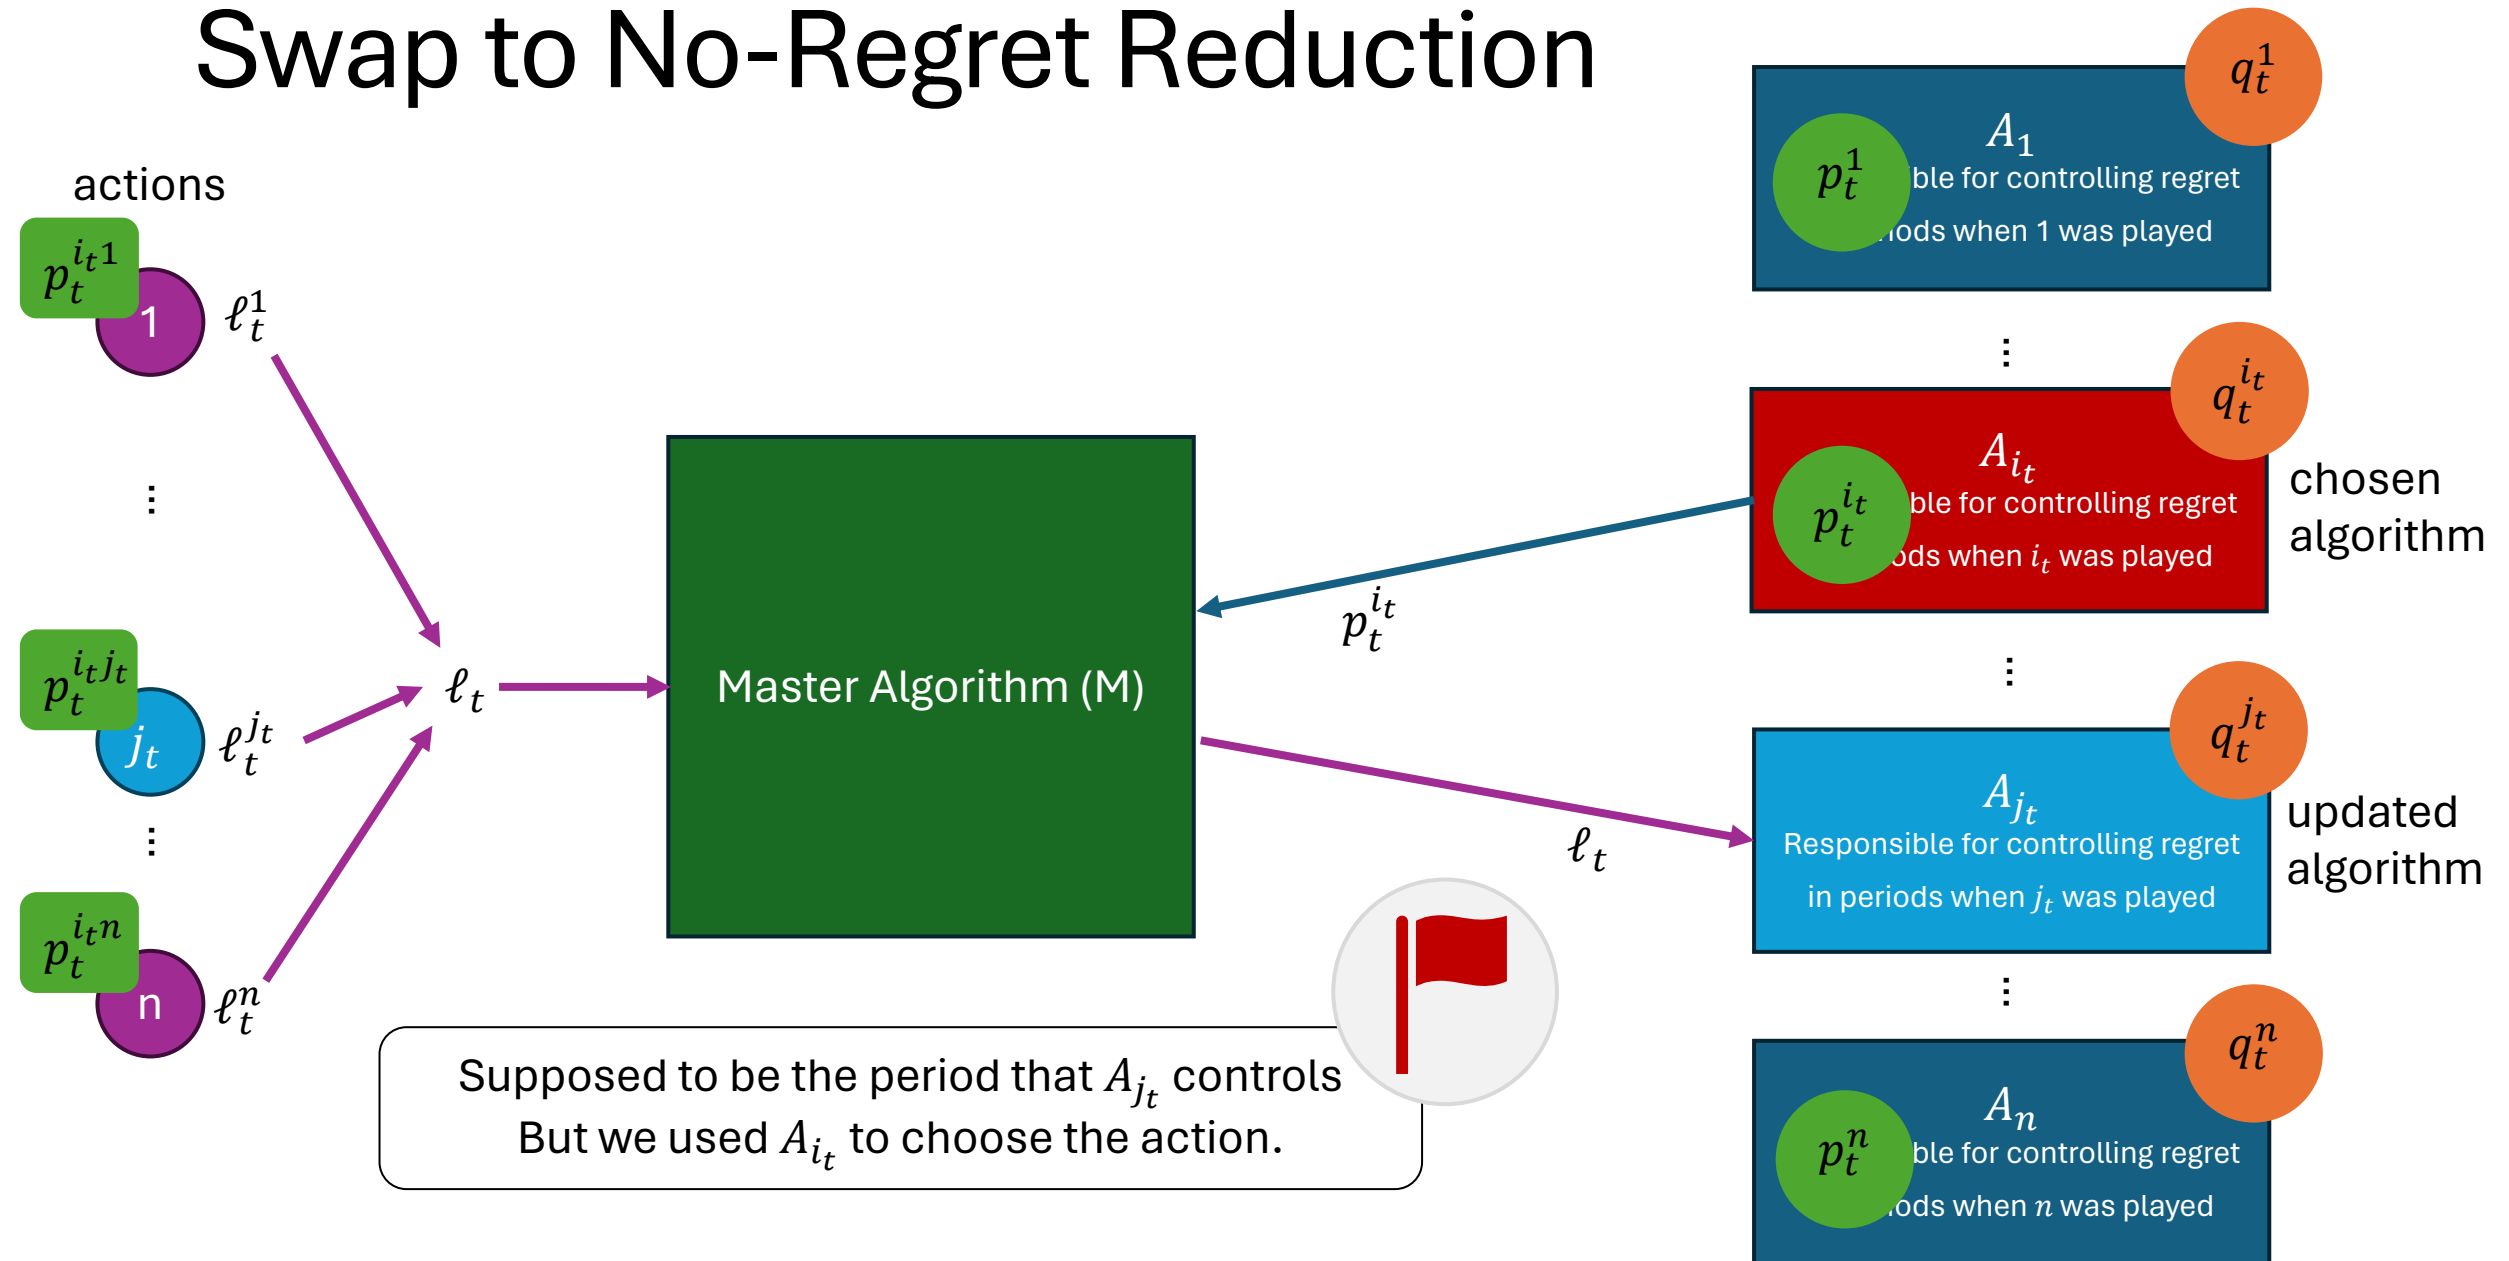
Swap to No-Regret Reduction



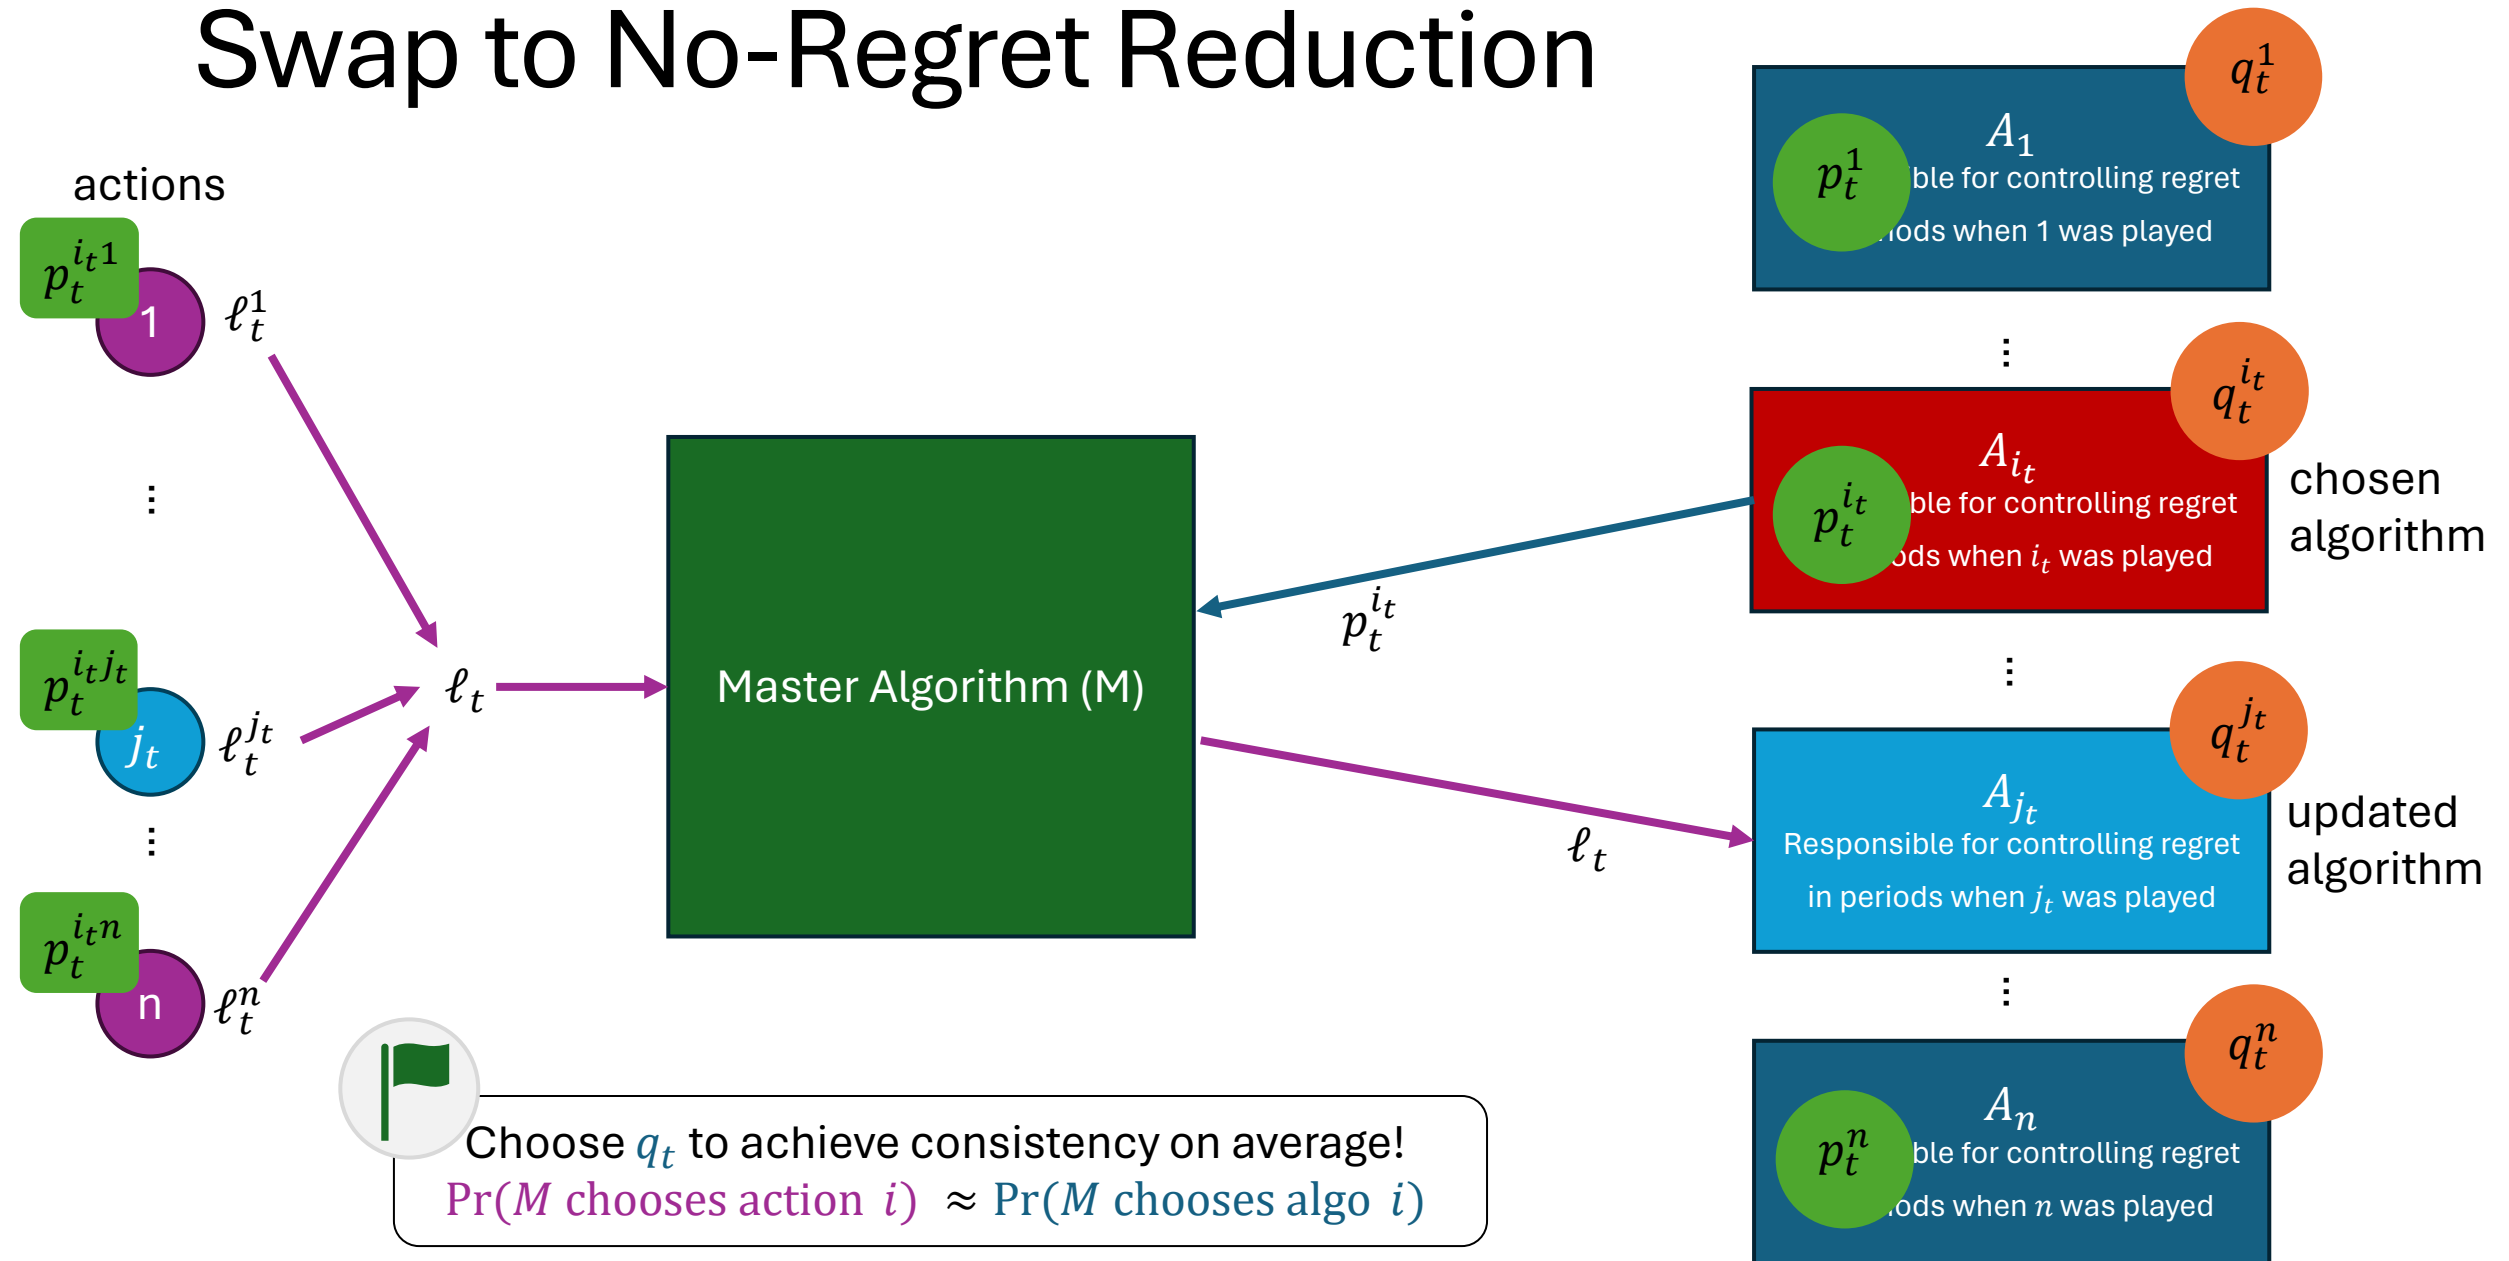
Swap to No-Regret Reduction



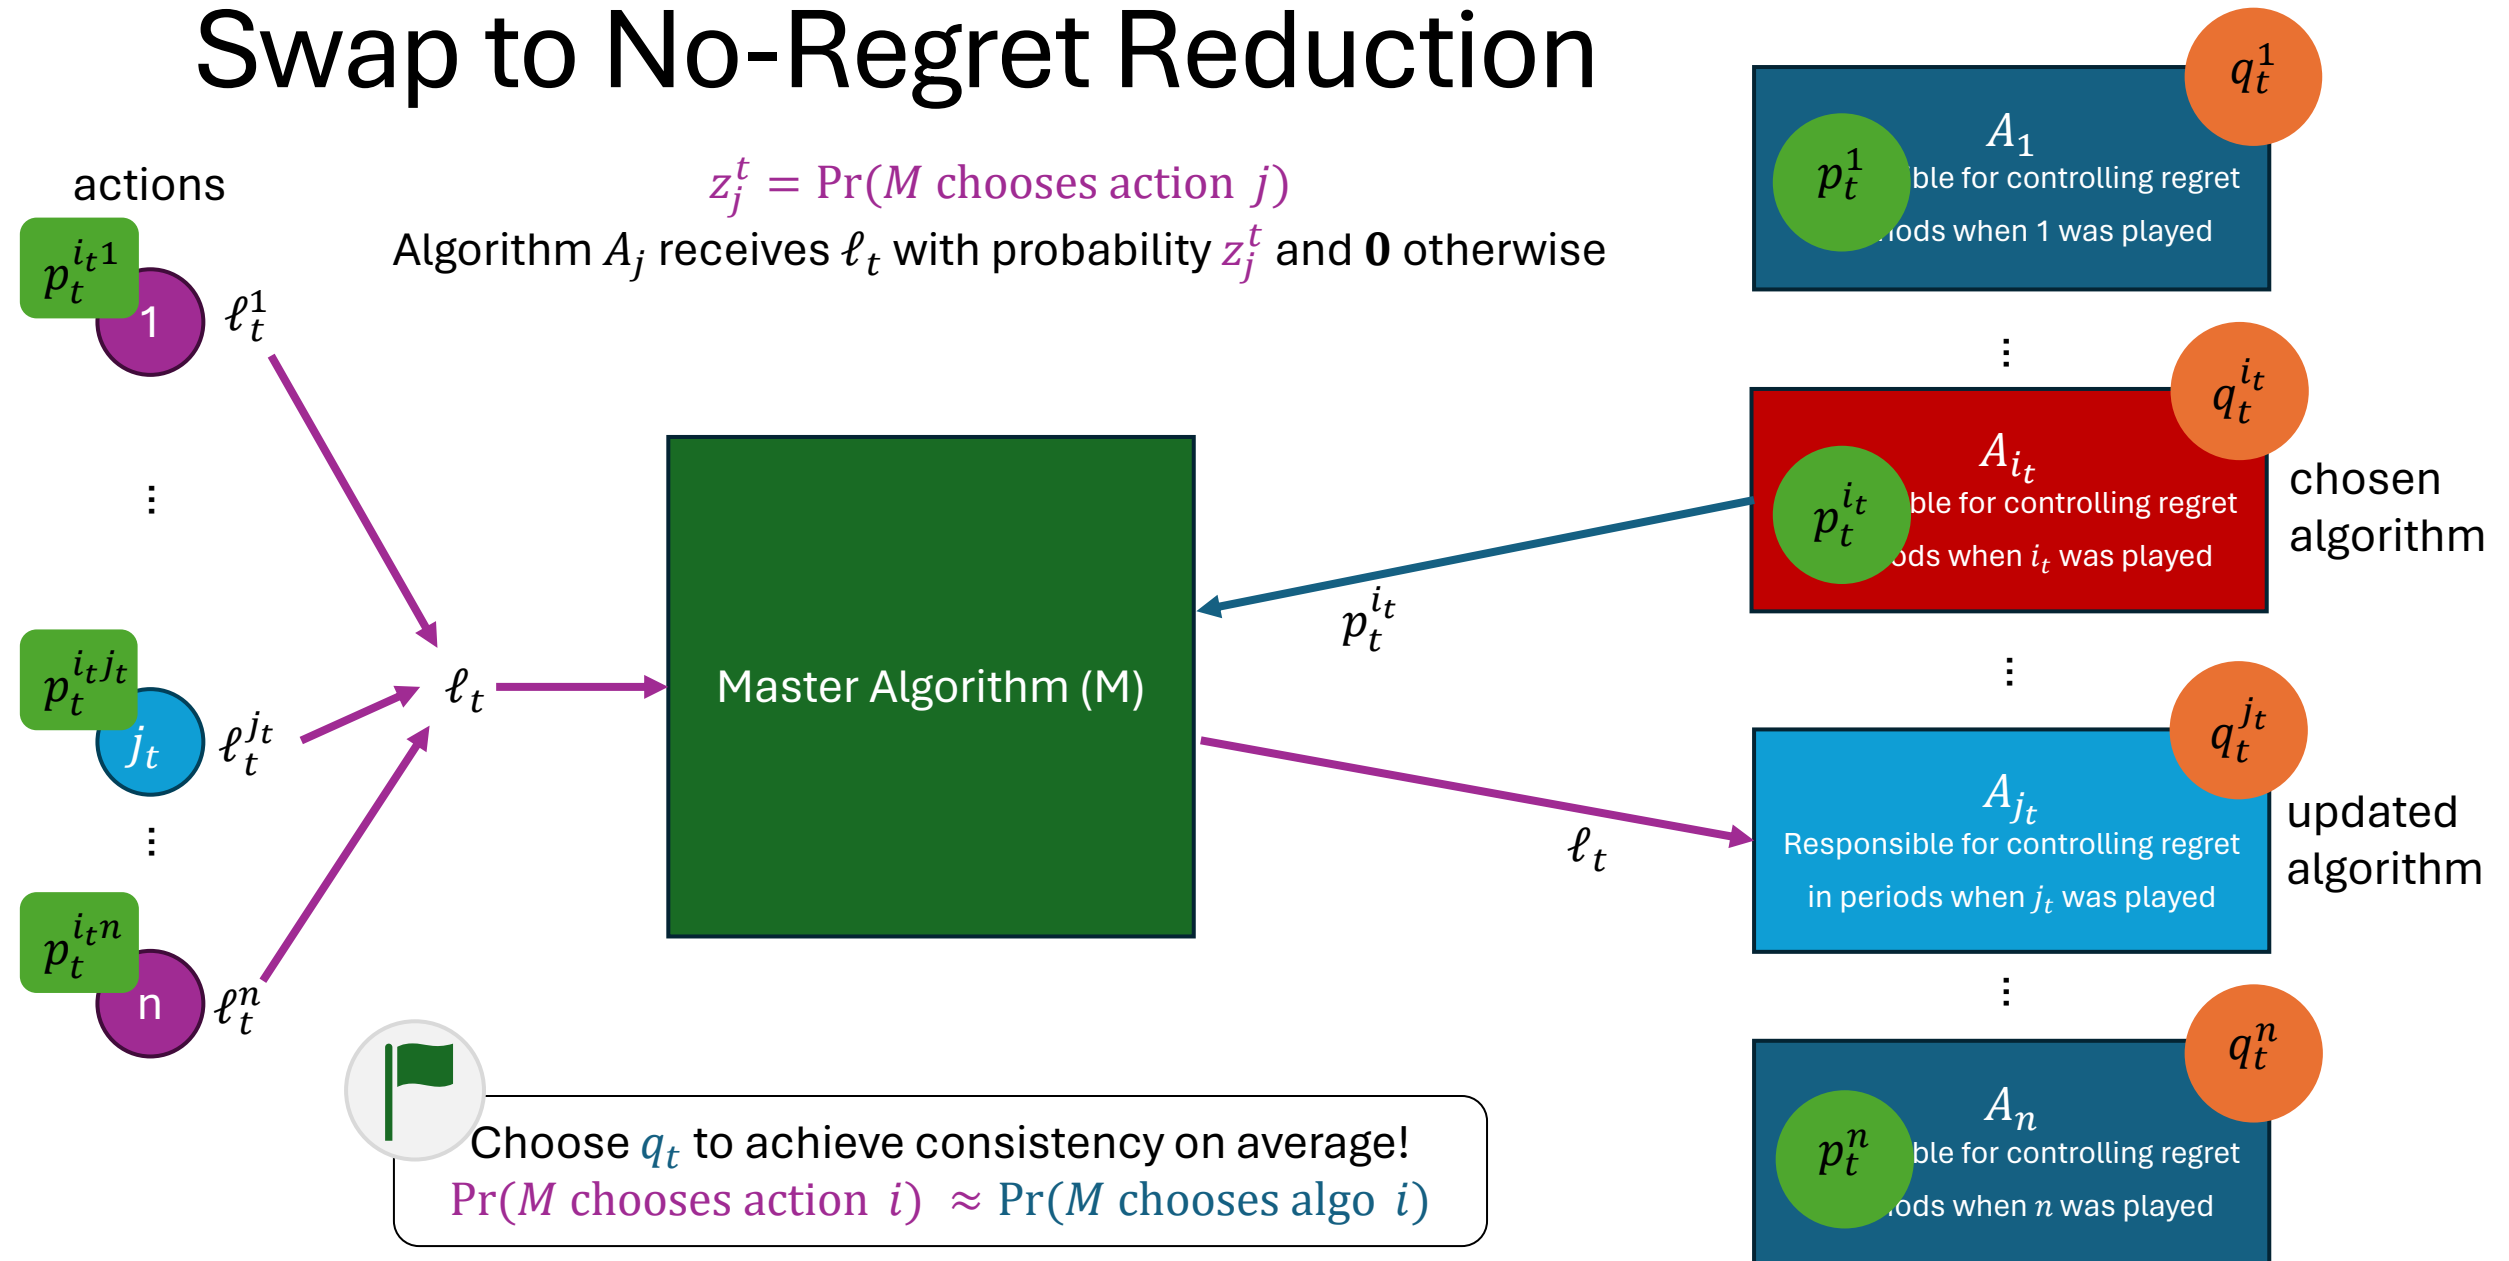
Swap to No-Regret Reduction



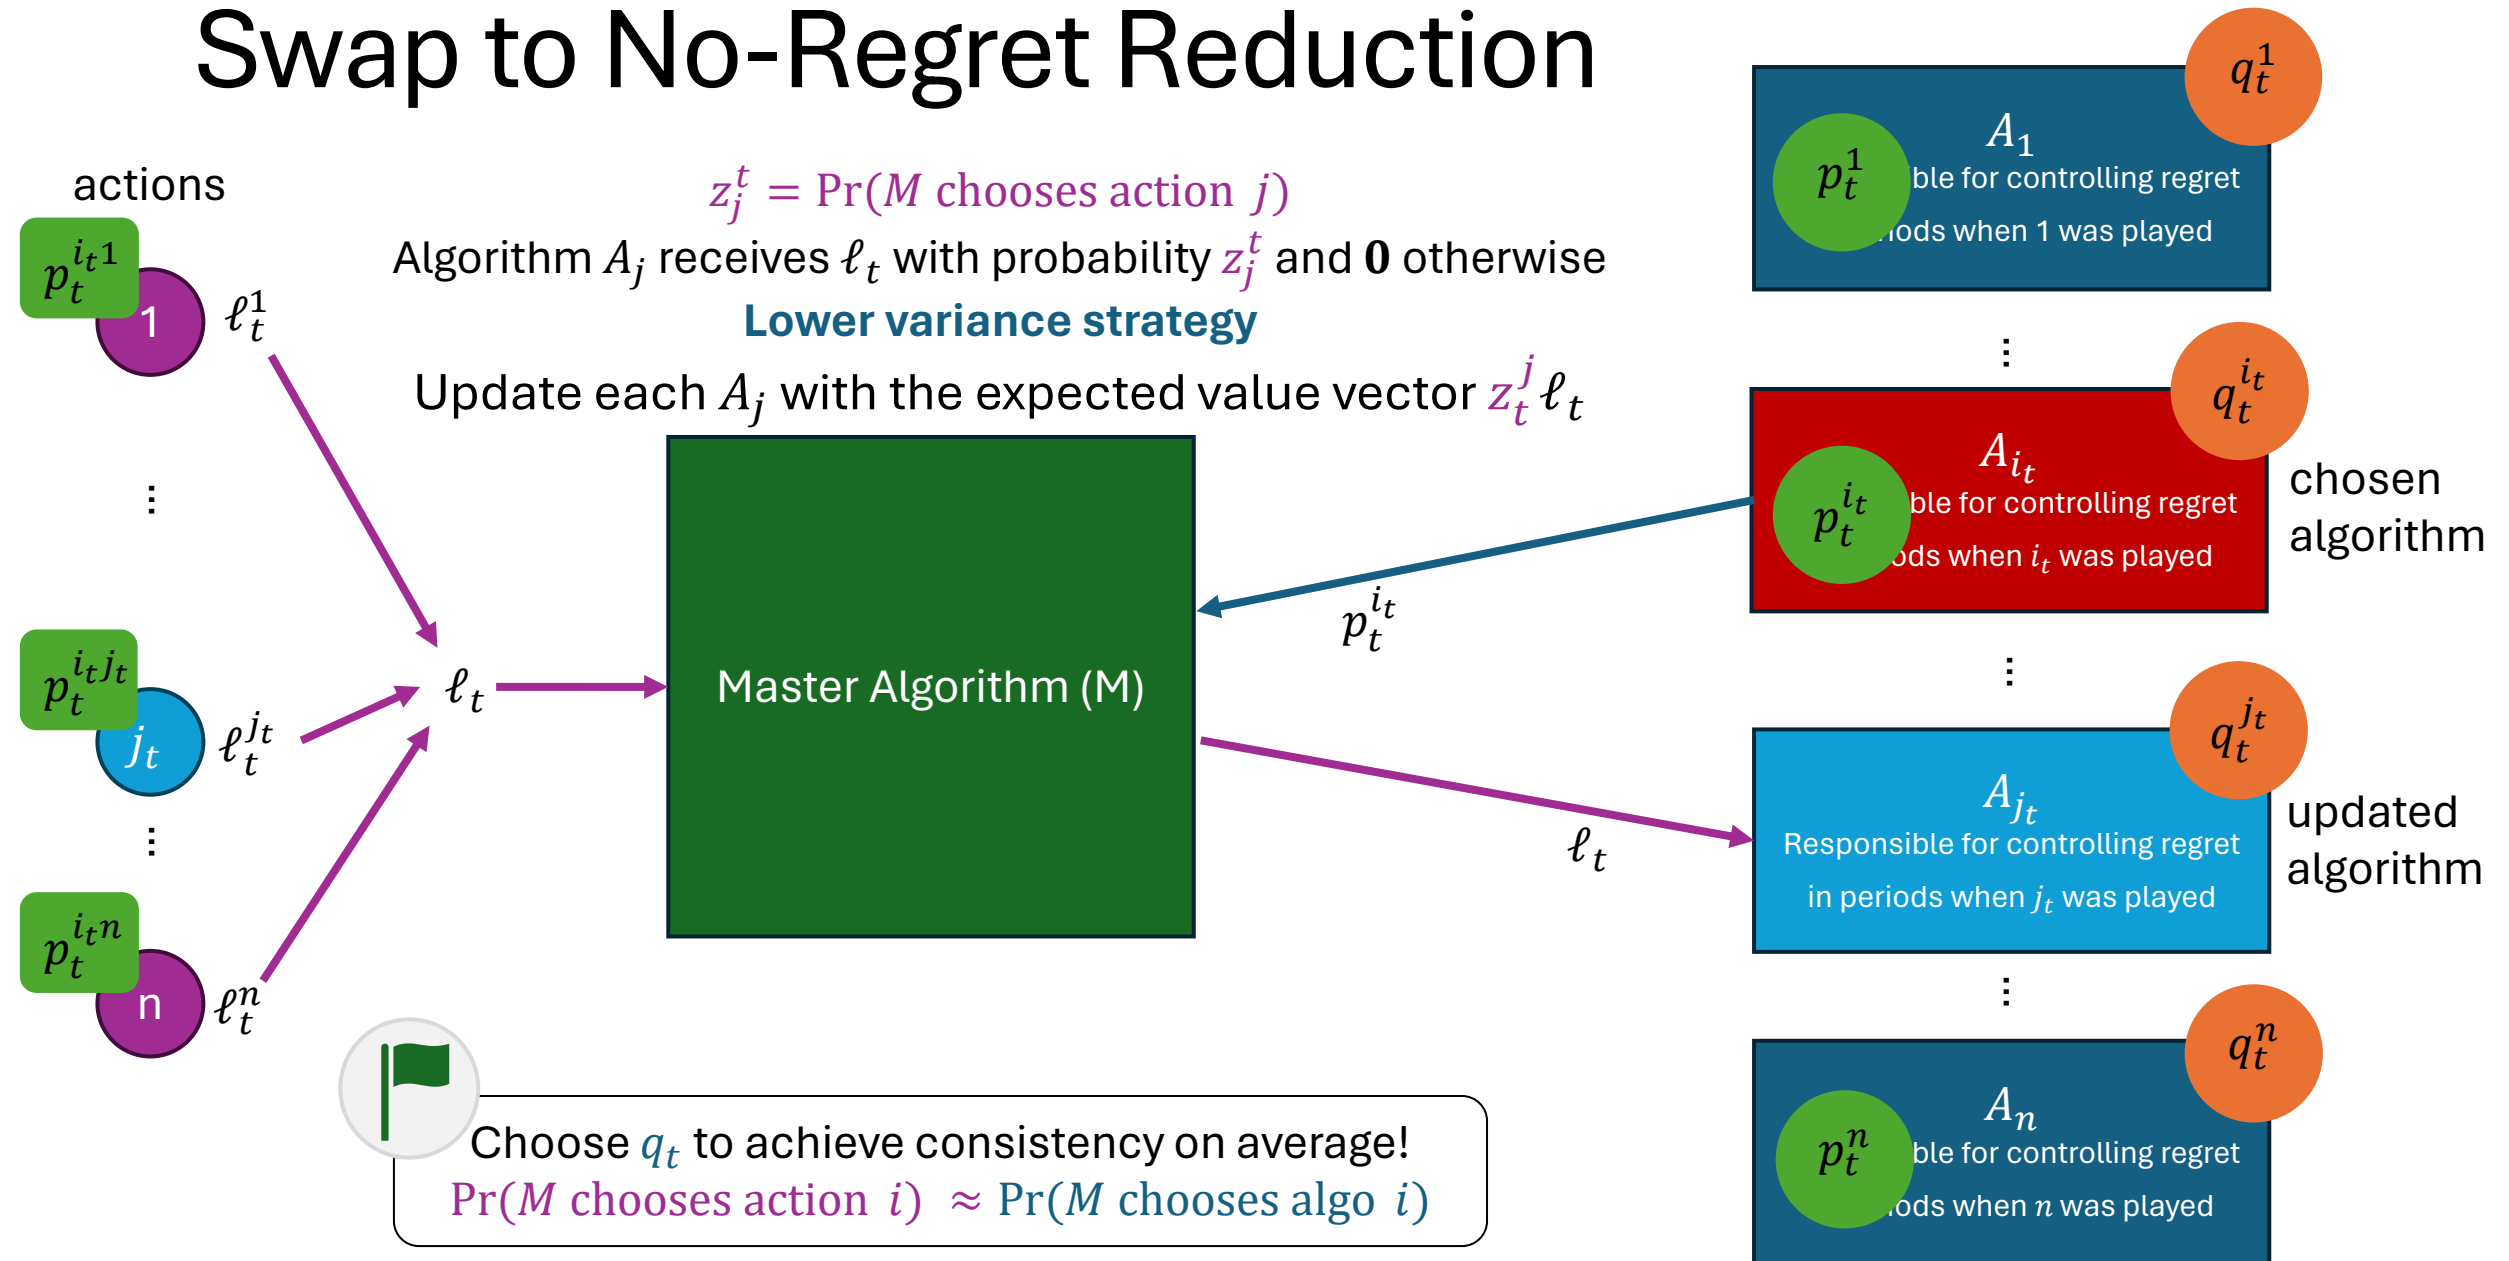
Swap to No-Regret Reduction



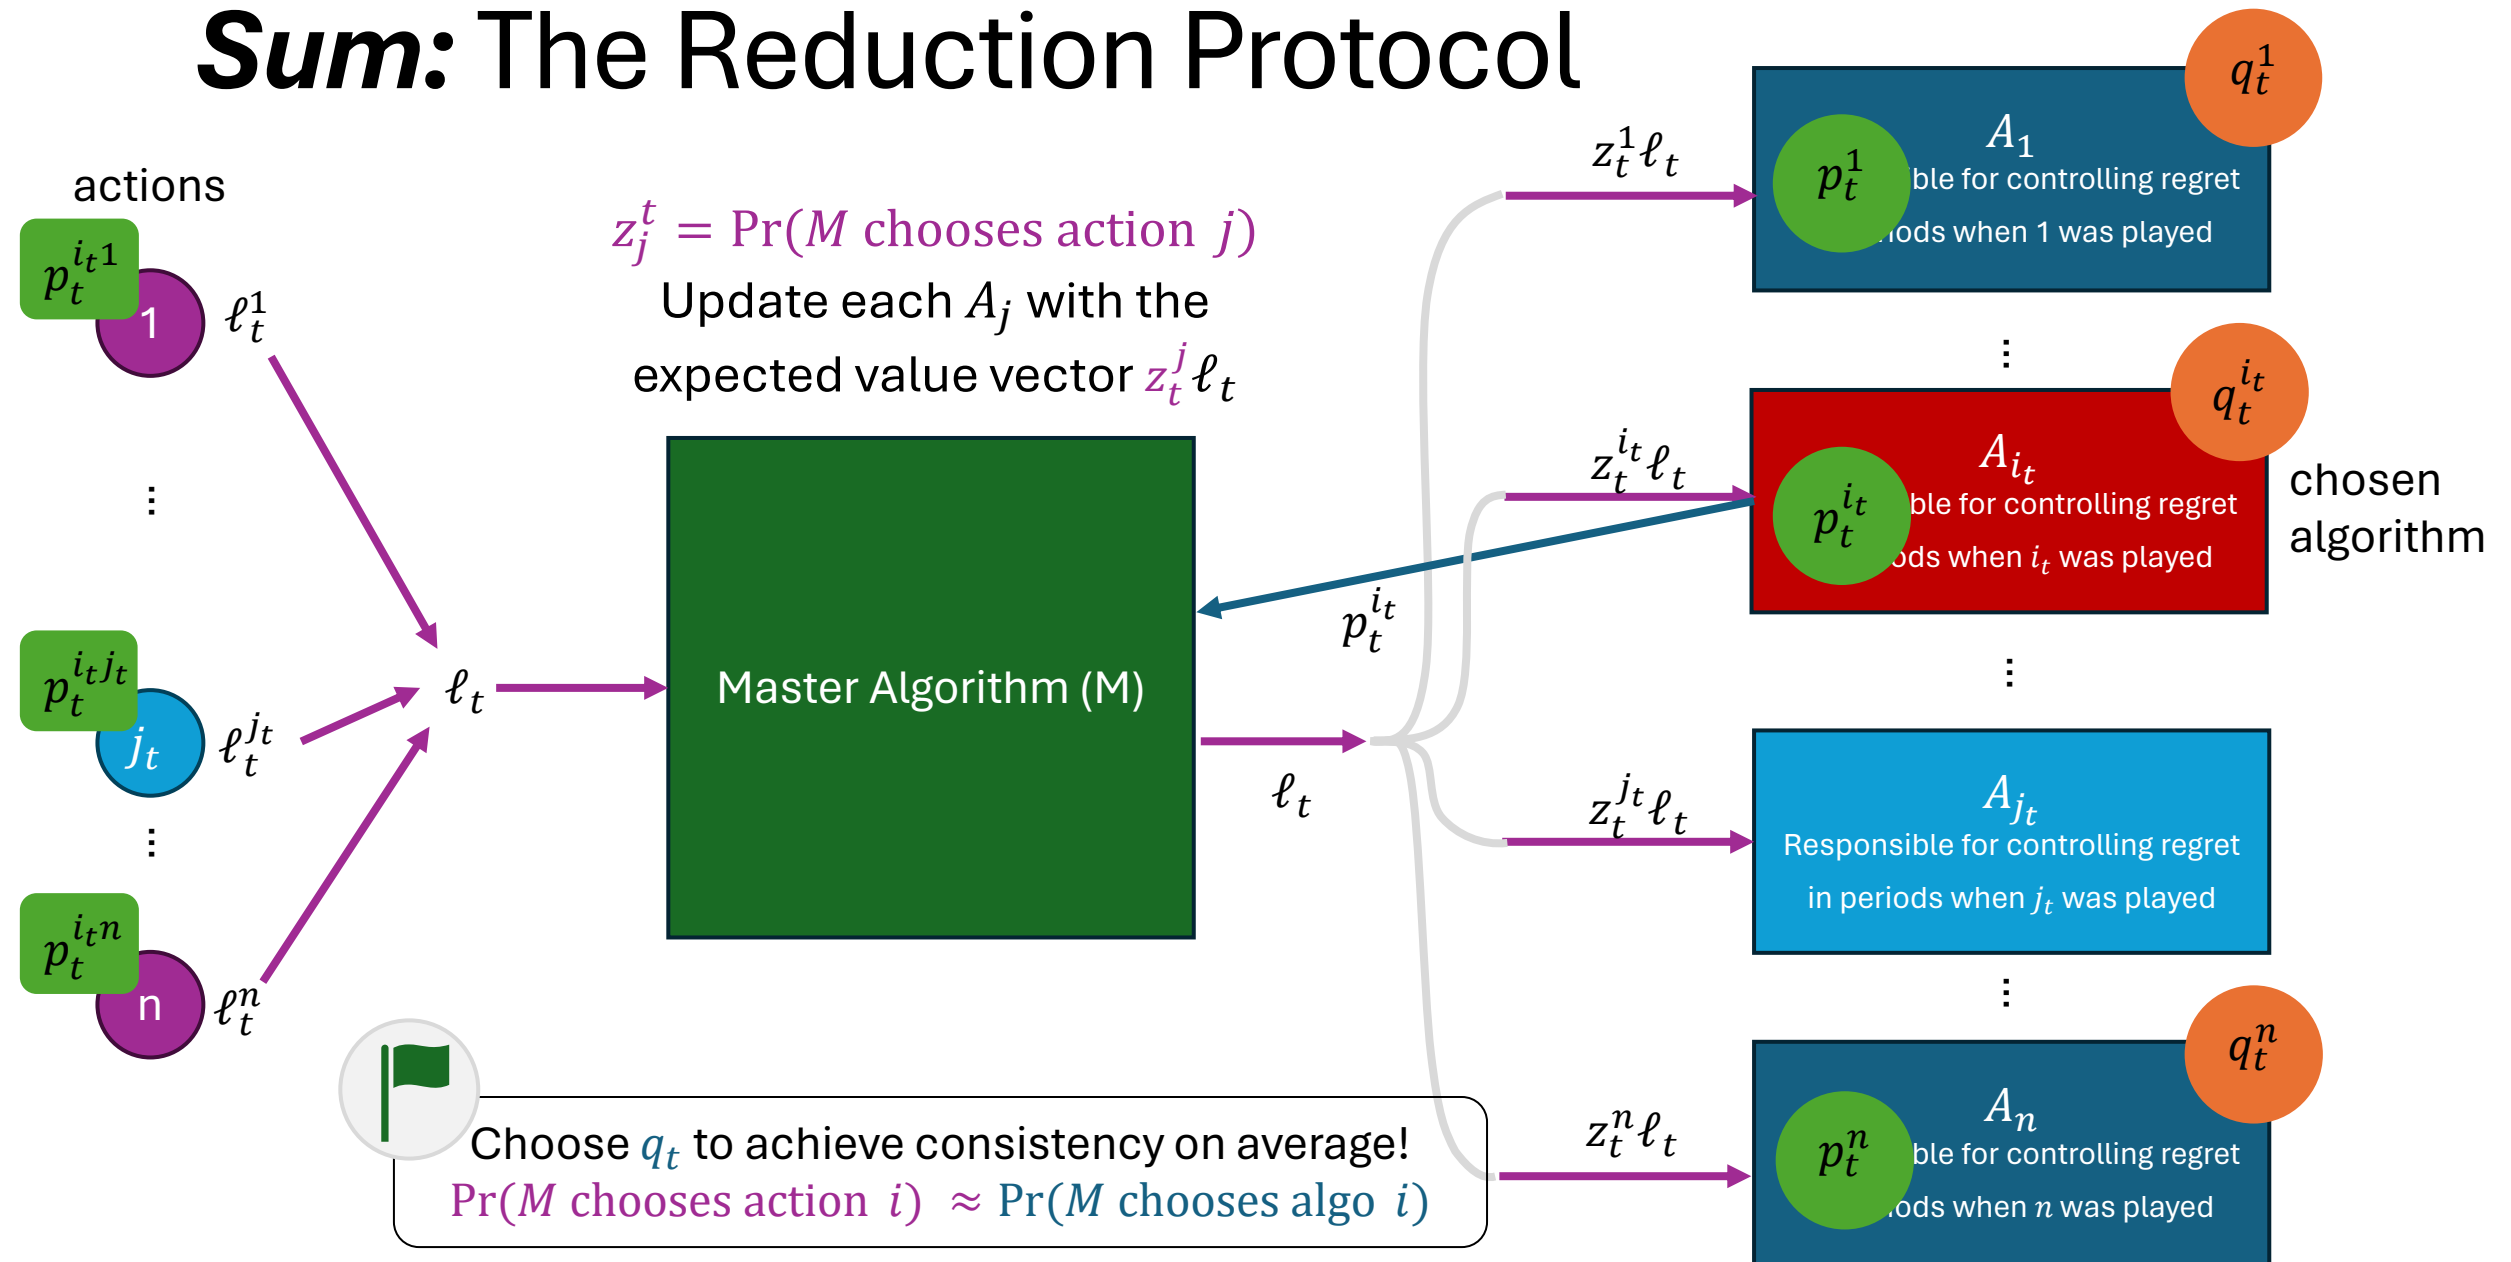
Swap to No-Regret Reduction



Swap to No-Regret Reduction



Sum: The Reduction Protocol



Sum: The reduction protocol

- At each period we choose each action with probability

$$\begin{aligned} z_t^j &= \Pr(M \text{ choose action } j) \\ &= \sum_i \underbrace{\Pr(M \text{ choose algo } A_i)}_{q_t^i} \cdot \underbrace{\Pr(A_i \text{ choose action } j)}_{p_t^{ij}} \end{aligned}$$

- We update each algorithm A_j with loss vector

$$z_t^j \ell_t = \Pr(M \text{ choose action } j) \cdot (\text{loss vector})$$

- The distribution over algorithms q_t is chosen such that

$$\Pr(M \text{ choose action } j) \approx \Pr(M \text{ choose algo } A_j)$$

From No-Regret of Algos
to No-Swap Regret of Master

$$\text{Regret} = \text{Loss} - \text{Benchmark Loss}$$

Loss Analysis at Each Step

- How much loss does algorithm A_i perceive?

$$\frac{\text{Pr}(M \text{ choose action } i)}{\text{The fraction of the loss vector that } M \text{ attributed and reported back to } A_i} \sum_j \text{Pr}(A_i \text{ choose action } j) \cdot \text{loss}(j)$$

- How much total loss do all the algorithms perceive?

$$\sum_i \text{Pr}(M \text{ choose action } i) \sum_j \text{Pr}(A_i \text{ choose action } j) \cdot \text{loss}(j)$$

- How much loss does the master algorithm incur?

$$\sum_i \text{Pr}(M \text{ choose algo } A_i) \sum_j \text{Pr}(A_i \text{ choose action } j) \cdot \text{loss}(j)$$

Loss Analysis at Each Step

- How much loss does algorithm A_i perceive?

$$\frac{\text{Pr}(M \text{ choose action } i)}{\text{The fraction of the loss vector that } M \text{ attributed and reported back to } A_i} \sum_j \text{Pr}(A_i \text{ choose action } j) \cdot \text{loss}(j)$$

- How much total loss do all the algorithms perceive?

$$\sum_i \text{Pr}(M \text{ choose action } i) \sum_j \text{Pr}(A_i \text{ choose action } j) \cdot \text{loss}(j)$$

- How much loss does the master algorithm incur?

$$\sum_i \text{Pr}(M \text{ choose algo } A_i) \sum_j \text{Pr}(A_i \text{ choose action } j) \cdot \text{loss}(j)$$

Recap: Loss Analysis at Each Step

Corollary. If we can guarantee that

$$\underbrace{\Pr(M \text{ choose action } i)}_{z_t^i} \approx \underbrace{\Pr(M \text{ choose algo } A_i)}_{q_t^i}$$

Then the total loss perceived by the separate algorithms is approximately the same as the total loss experienced by the master

$$\text{total loss perceived by algos} \approx \text{total loss of master}$$

Competing Benchmark Analysis at Each Step

- What can each algorithm A_i compete with based on **no-regret**?

$$\Pr(M \text{ choose action } i) \cdot \text{loss}(\phi(i))$$

The fraction of the loss vector that M attributed and reported back to A_i

For each algo A_i this is a constant action comparison with $i' = \phi(i)$

- What can in total all algorithms compete with based on **no-regret**?

$$\sum_i \Pr(M \text{ choose action } i) \cdot \text{loss}(\phi(i))$$

- What does the master want to compete with for **no-swap regret**?

$$\sum_j \Pr(M \text{ choose action } j) \cdot \text{loss}(\phi(j))$$

Competing Benchmark Analysis at Each Step

- What can each algorithm A_i compete with based on **no-regret**?

$$\Pr(M \text{ choose action } i) \cdot \text{loss}(\phi(i))$$

The fraction of the loss vector that M attributed and reported back to A_i

For each algo A_i this is a constant action comparison with $i' = \phi(i)$

- What can in total all algorithms compete with based on **no-regret**?

$$\sum_i \Pr(M \text{ choose action } i) \cdot \text{loss}(\phi(i))$$

- What does the master want to compete with for **no-swap regret**?

$$\sum_j \Pr(M \text{ choose action } j) \cdot \text{loss}(\phi(j))$$

Recap: Benchmark Analysis at Each Step

Corollary. The total *perceived* benchmark loss that algorithms compete with, where each algorithm i considers the no-regret benchmark of always playing action $i' = \phi(i)$, is equal to the *true* swap benchmark loss that the master wants to compete with, associated with the swap function ϕ .

$$\text{Regret} = \text{Loss} - \text{Benchmark Loss}$$

Regret Analysis at Each Step

Corollary. If we can guarantee that

$$\Pr(M \text{ choose action } i) \approx \Pr(M \text{ choose algo } A_i)$$

then swap regret of master is upper bounded by sum of plain regrets of algos

$$\text{Swap Regret of Master} = \text{Total Loss of Master} - \text{Swap Benchmark}$$

$$\approx \text{Total Perceived Loss by Algos} - \text{Total Algo Fixed Action Benchmark}$$

$$= \text{Total Perceived Regret of Algos}$$

Regret Analysis at Each Step

Corollary. If we can guarantee that

$$\Pr(M \text{ choose action } i) \approx \Pr(M \text{ choose algo } A_i)$$

then swap regret of master is upper bounded by sum of plain regrets of algos

$$\begin{aligned} \sum_t \sum_j z_t^j \ell_t^j - z_t^j \ell_t^{\phi(j)} &= \sum_t \sum_i q_t^i \sum_j p_t^{ij} \ell_t^j - \sum_t \sum_j z_t^j \ell_t^{\phi(j)} \\ &\approx \sum_t \sum_i z_t^i \sum_j p_t^{ij} \ell_t^j - \sum_t \sum_i z_t^i \ell_t^{\phi(i)} \\ &= \sum_i \sum_t \langle p_t^i, z_t^i \ell_t \rangle - z_t^i \ell_t^{\phi(i)} \end{aligned}$$

Can we pick q_t such that:

$$\Pr(M \text{ choose action } j) \approx \Pr(M \text{ choose algo } A_j)$$

Choosing distribution over algos

- Choose q_t such that

$$\Pr(M \text{ choose action } j) \approx \Pr(M \text{ choose algo } A_j)$$

- Remember that


$$\Pr(M \text{ choose action } j) = \sum_i \Pr(M \text{ choose algo } A_i) \cdot \Pr(A_i \text{ choose action } j)$$

- We need the distribution over algos q_t to satisfy the self-consistency property

$$\sum_i \underbrace{\Pr(M \text{ choose algo } A_i)}_{q_t^i} \cdot \underbrace{\Pr(A_i \text{ choose action } j)}_{p_t^{ij}} = \underbrace{\Pr(M \text{ choose algo } A_j)}_{q_t^j}$$

Does there exist a distribution q_t such that:

$$\sum_i \Pr(M \text{ choose algo } A_i) \cdot \Pr(A_i \text{ choose action } j) = \Pr(M \text{ choose algo } A_j)$$




A diagram consisting of two blue lines. The first line starts under the summation symbol \sum_i of the equation above and slopes downwards to the right, ending under the first q_t^i term of the equation below. The second line starts under the p_t^{ij} term of the equation below and slopes upwards to the right, ending under the A_j term of the equation above.

$$\sum_{i=1}^n q_t^i \cdot p_t^{ij} = q_t^j$$

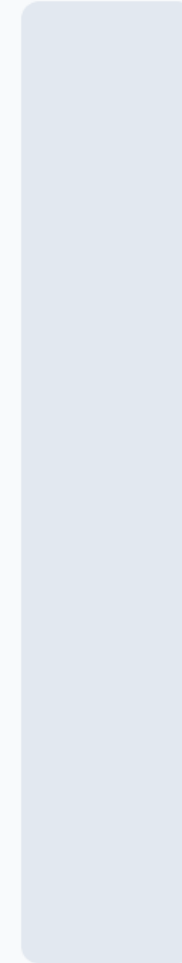
There always exists a distribution q that satisfies this property

0%



True


0%



False

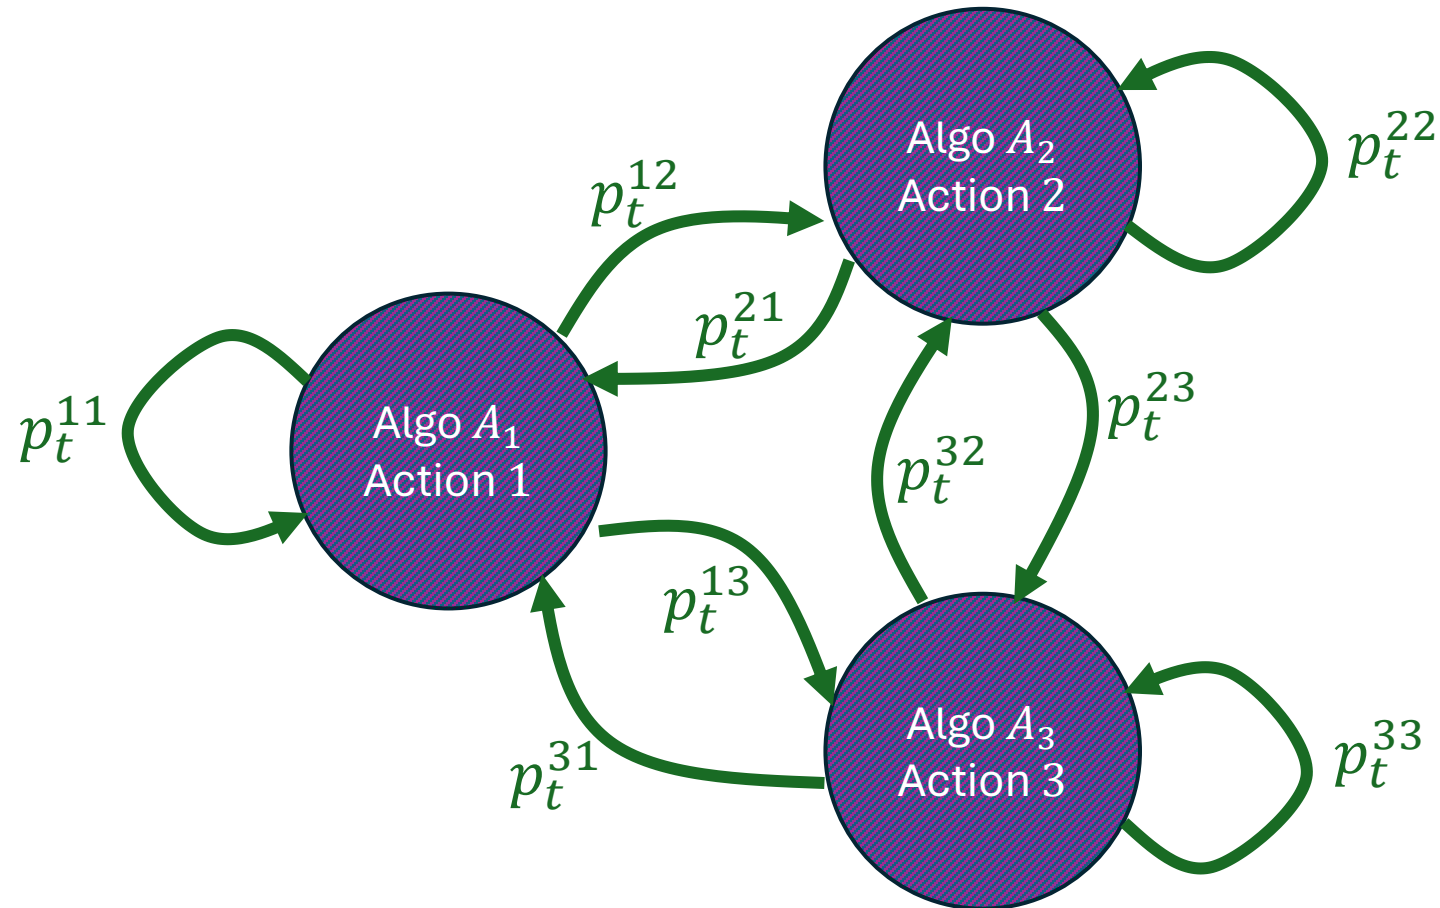
Choosing distribution over algos

$$\sum_i \Pr(M \text{ choose algo } A_i) \cdot \Pr(A_i \text{ choose action } j) = \Pr(M \text{ choose algo } A_j)$$



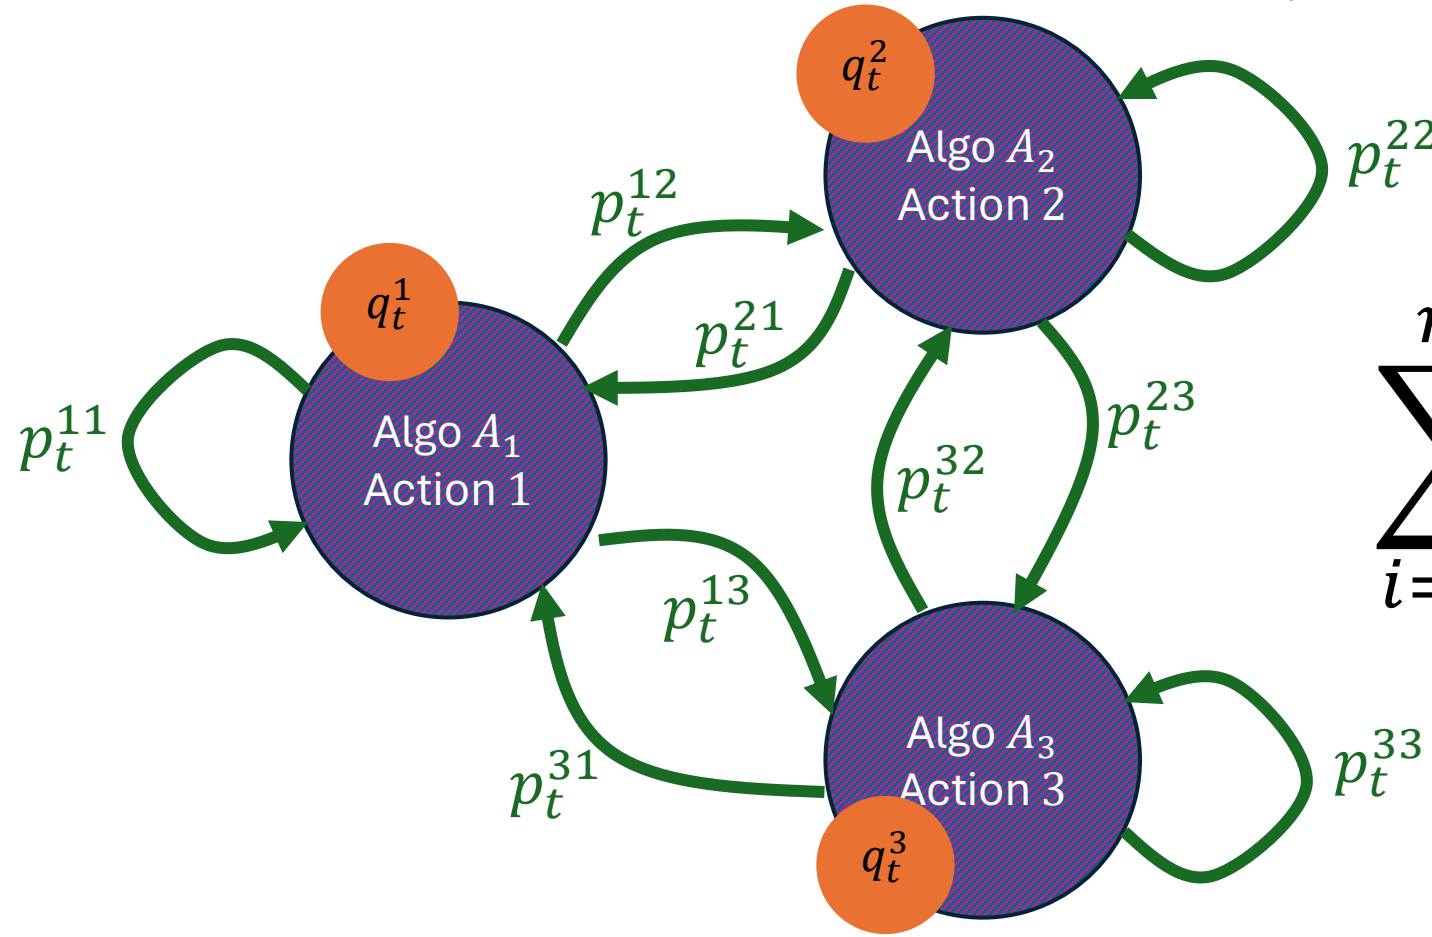
Choosing distribution over algos

$$\sum_i \Pr(M \text{ choose algo } A_i) \cdot \Pr(A_i \text{ choose action } j) = \Pr(M \text{ choose algo } A_j)$$



Choosing distribution over algos


$$\sum_i \Pr(M \text{ choose algo } A_i) \cdot \Pr(A_i \text{ choose action } j) = \Pr(M \text{ choose algo } A_j)$$



$$\sum_{i=1}^n q_t^i \cdot p_t^{ij} = q_t^j$$

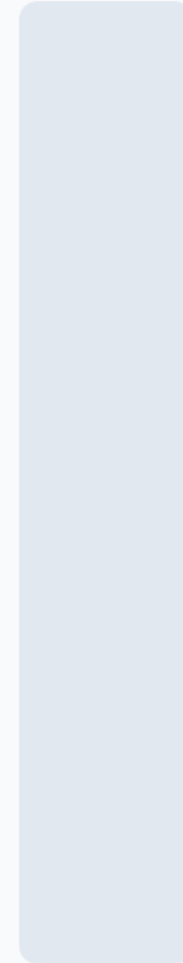
There always exists a distribution q that satisfies this property

0%



True

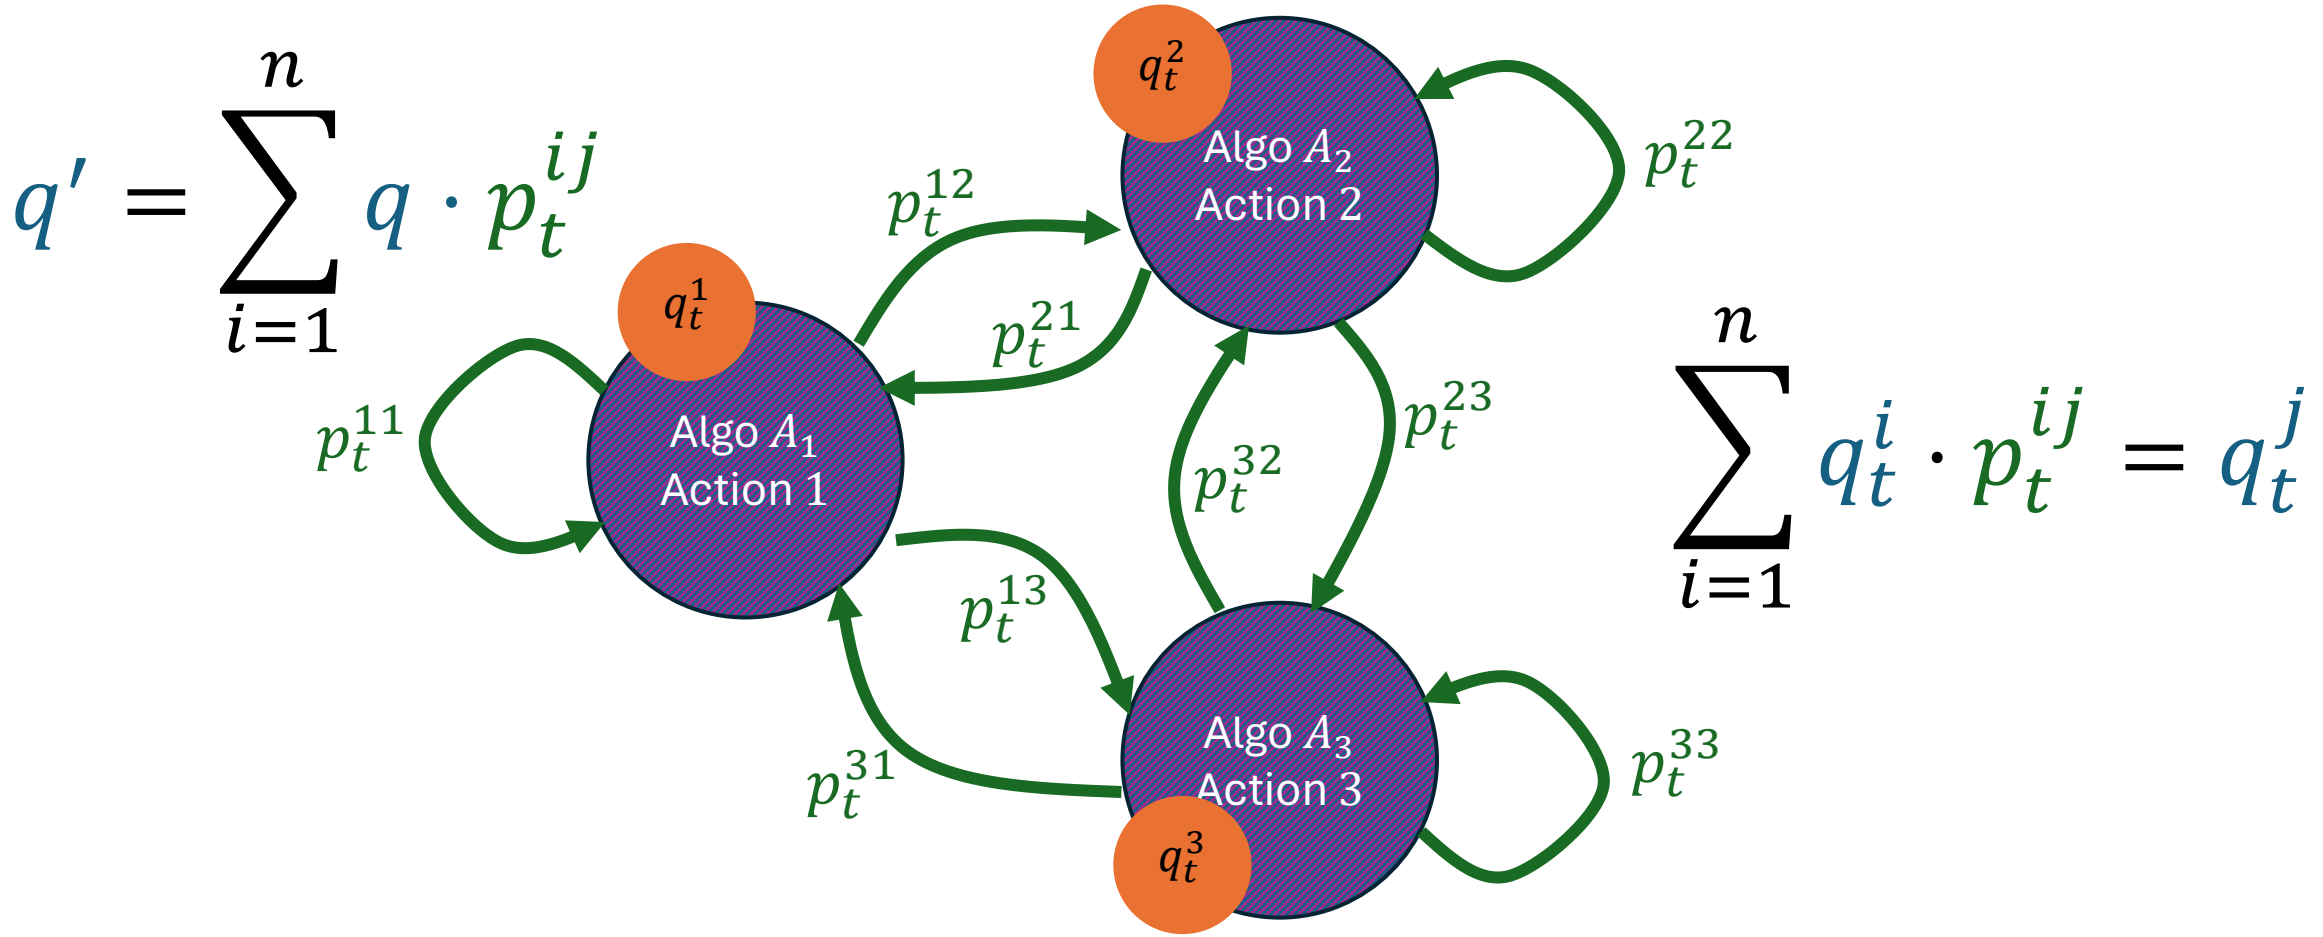
0%



False

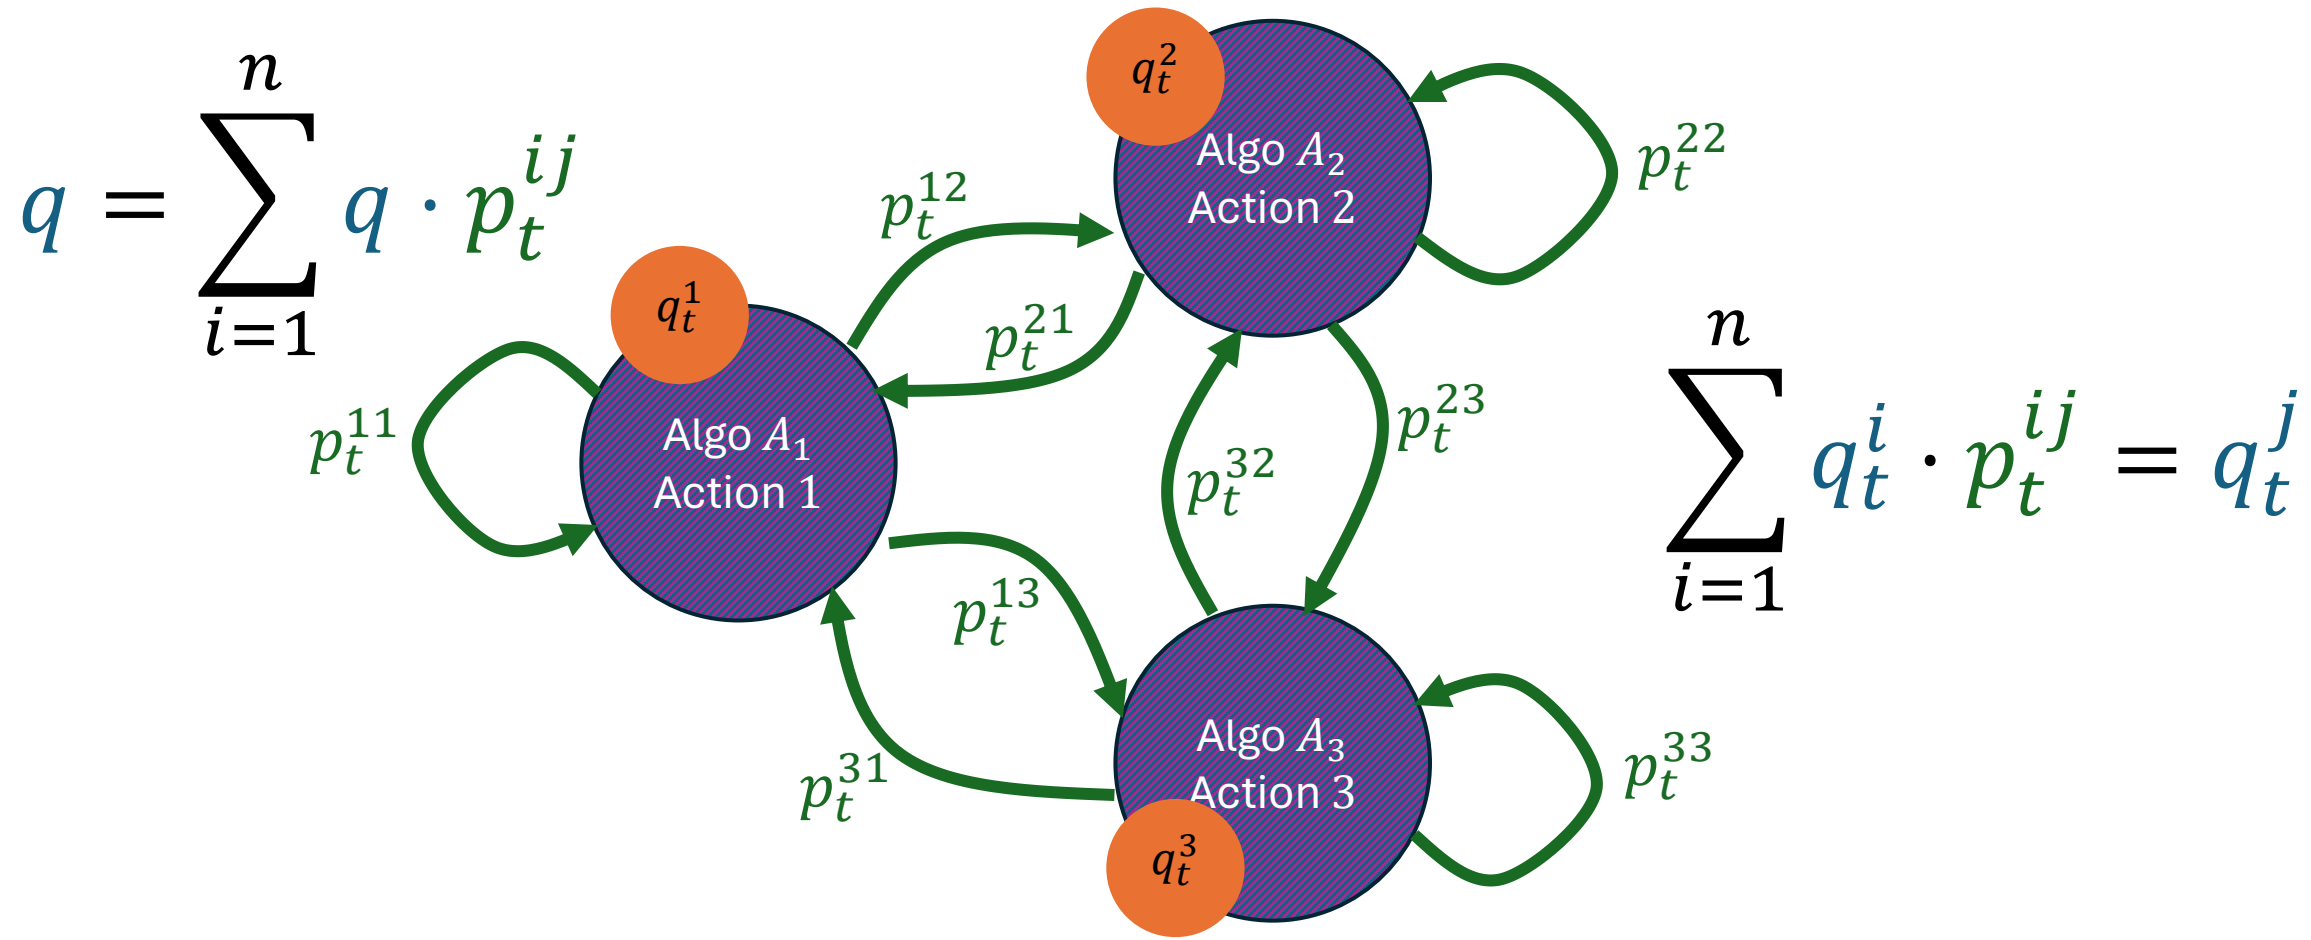
A Markov Chain over the Algos/Actions

Starting from a distribution q over nodes and applying one step of the random transitions, brings us to a new distribution over states



Stationary Distributions of Markov Chains

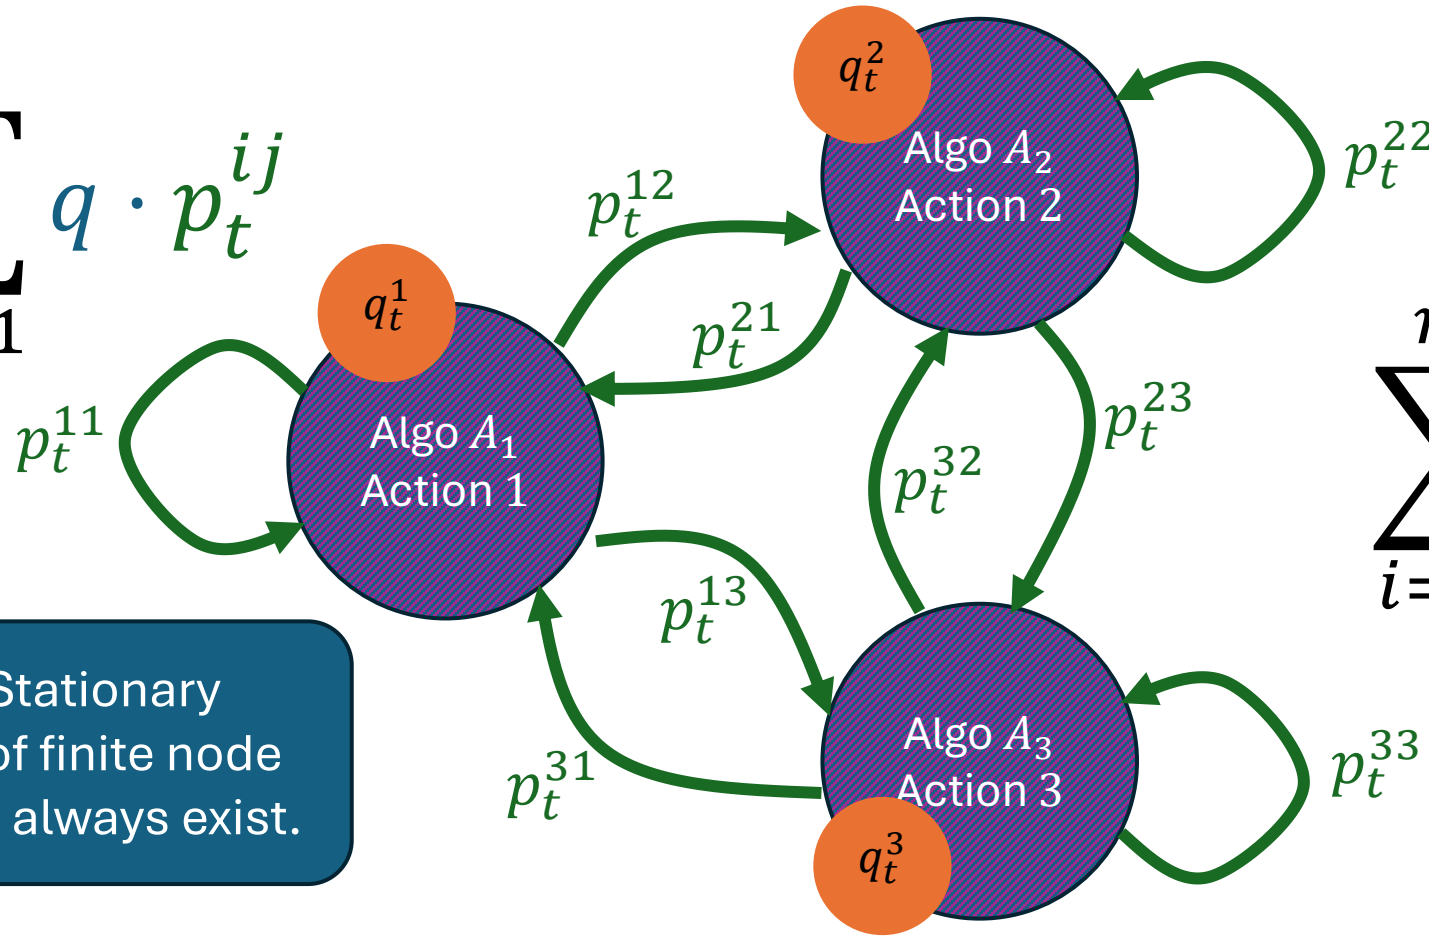
If new distribution is the same as the original distribution, then this distribution is called a **Stationary Distribution of the Markov Chain**



Stationary Distributions of Markov Chains

If new distribution is the same as the original distribution, then this distribution is called a **Stationary Distribution of the Markov Chain**

$$q = \sum_{i=1}^n q \cdot p_t^{ij}$$



$$\sum_{i=1}^n q_t^i \cdot p_t^{ij} = q_t^j$$

Theorem. Stationary distributions of finite node Markov Chains always exist.

Recap: Choosing Distribution over Algos

Corollary. If we choose q_t as stationary distribution of the Markov Chain defined by transition probabilities $\Pr(i \rightarrow j) = p_t^{ij}$ then

$$\Pr(M \text{ choose action } j) = \Pr(M \text{ choose algo } A_j)$$

Therefore

$$\text{Swap Regret of Master} = \text{Total Fixed Action Regret of Algos} \rightarrow 0$$

Sum: The reduction protocol

- At each period calculate stationary distribution q_t of the Markov Chain defined by the transition probabilities $\Pr(i \rightarrow j) = p_t^{ij}$
- Choose each action with probability

$$z_t^j = \Pr(M \text{ choose action } j) = \Pr(M \text{ choose algo } j) = q_t^j$$

- Update each algorithm A_j with loss vector

$$z_t^j \ell_t = \Pr(M \text{ choose action } j) \cdot (\text{loss vector})$$

Finding Stationary Distributions

- Define the matrix P_t , whose (i, j) entry is p_t^{ij}
- Then the stationary distribution satisfies
$$q^\top = q^\top P_t$$
- q is a left eigenvector of P_t associated with eigenvalue 1
- We can calculate q via eigen-decomposition of P_t and identifying the eigenvector associated with eigenvalue 1

Overall Algorithm using EXP for each Algo

```
Initialize Pt with each row being the uniform distribution
For t in 1..T
    # Calculate choice probability q of master based on
    # choice probabilities Pt of algos
    Calculate stationary distribution q of matrix Pt
    Draw action jt based on distribution q
    Observe loss vector lt

    # update each algorithms choice probabilities
    For i in 1..n
        Calculate perceived loss plt[i] = q[i] * lt
        Pt[i] = EXP-Update(Pt[i], plt[i])
```

Recap: Final Theorem

Theorem. If we choose q_t as stationary distribution of the Markov Chain defined by transition probabilities $\Pr(i \rightarrow j) = p_t^{ij}$ and each algorithm updates their choice probabilities using the EXP rule then

$$\text{Average Swap Regret of Master} \leq 2n \sqrt{\frac{2 \log(n)}{T}} \rightarrow 0$$

Convergence to Correlated Equilibrium

Theorem. If all players use such an algorithm, then the empirical joint distribution of actions converges to the set of correlated equilibria.

At every T the empirical joint distribution of strategies π^T is an $\epsilon(T)$ approximate correlated equilibrium, in the sense that:

$$\text{SwapRegret}_i(s_i, s'_i, T) = \sum_{s_{-i}} \pi^T(s_i, s_{-i}) \cdot \left(u_i(s_i, s_{-i}) - u_i(s'_i, s_{-i}) \right) \leq \epsilon(T)$$

with $\epsilon(T) = 2n \sqrt{\frac{2 \log(n)}{T}}$, where n is number of actions of player i

Recent example research in multi-agent RL using Correlated Equilibrium Techniques

Multi-Agent Training beyond Zero-Sum with Correlated Equilibrium Meta-Solvers

Luke Marris^{1,2} Paul Muller^{1,3} Marc Lanctot¹ Karl Tuyls¹ Thore Graepel^{1,2}

Abstract

Two-player, constant-sum games are well studied in the literature, but there has been limited progress outside of this setting. We propose Joint Policy-Space Response Oracles (JPSRO), an algorithm for training agents in n-player, general-sum extensive form games, which provably converges to an equilibrium. We further suggest correlated equilibria (CE) as promising meta-solvers, and propose a novel solution concept Maximum Gini Correlated Equilibrium (MGCE), a principled and computationally efficient family of solutions for solving the correlated equilibrium selection problem. We conduct several experiments using CE meta-solvers for JPSRO and demonstrate convergence on n-player, general-sum games.

1. Introduction

Recent success in tackling two-player, constant-sum games (Silver et al., 2016; Vinyals et al., 2019) has outpaced progress in n-player, general-sum games despite a lot of interest (Jaderberg et al., 2019; OpenAI et al., 2019; Brown & Sandholm, 2019; Lockhart et al., 2020; Gray et al., 2020; Anthony et al., 2020). One reason is because Nash equilibrium (NE) (Nash, 1951) is tractable and interchangeable in the two-player, constant-sum setting but becomes intractable (Daskalakis et al., 2009) and potentially non-interchangeable¹ in n-player and general-sum settings. The problem of selecting from multiple solutions is known as the equilibrium selection problem (Goldberg et al., 2013;

Avis et al., 2010; Harsanyi & Selten, 1988).²

Outside of normal form (NF) games, this problem setting arises in multi-agent training when dealing with empirical games (also called meta-games), where a game payoff tensor is populated with expected outcomes between agents playing an extensive form (EF) game, for example the StarCraft League (Vinyals et al., 2019) and Policy-Space Response Oracles (PSRO) (Lanctot et al., 2017), a recent variant of which reached state-of-the-art results in Stratego Barrage (McAleer et al., 2020).

In this work we propose using correlated equilibrium (CE) (Aumann, 1974) and coarse correlated equilibrium (CCE) as a suitable target equilibrium space for n-player, general-sum games³. The (C)CE solution concept has two main benefits over NE; firstly, it provides a mechanism for players to correlate their actions to arrive at mutually higher payoffs and secondly, it is computationally tractable to compute solutions for n-player, general-sum games (Daskalakis et al., 2009). We provide a tractable approach to select from the space of (C)CEs (MG), and a novel training framework that converges to this solution (JPSRO). The result is a set of tools for theoretically solving any complete information⁴ multi-agent problem. These tools are amenable to scaling approaches; including utilizing reinforcement learning, function approximation, and online solution solvers, however we leave this to future work.

In Section 2 we provide background on a) correlated equilibrium (CE), an important generalization of NE, b) coarse correlated equilibrium (CCE) (Moulin & Vial, 1978), a similar solution concept, and c) PSRO, a powerful multi-agent training algorithm. In Section 3 we propose novel solution concepts called Maximum Gini (Coarse) Correlated Equilibrium (MG(C)CE) and in Section 4 we thoroughly explore its properties including tractability, scalability, invariance, and

¹DeepMind ²University College London ³Université Gustave Eiffel. Correspondence to: Luke Marris <marris@google.com>.

²The equilibrium selection problem is subtle and can have various interpretations. We describe it fully in Section 4.1 based